

Securing the Cloud: A Guide to Effective Vulnerability Management



Table of Contents

01	Intro	03
02	Vulnerability Management Then and Now	04
	The Framework.....	05
	Vulnerability Management for Software You Build	06
	Shift left	06
	Full Lifecycle Vulnerability Management.....	06
	Where are the vulnerabilities?	08
	When are the vulnerabilities discovered?	13
	Prioritization of Vulnerability Assessment Findings	15
	Runtime risk intelligence for cloud-native applications.....	16
	Remediation.....	17
	Mitigation	18
	Validation	20
	Risk acceptance.....	20
	Reporting	20
	Nuances of Cloud-Native Architecture	21
	Integrating Into the Development Lifecycle	22
	Binaries, packages, and language intricacies	24
	Update in flight or destroy and redeploy	26
	Using admission controllers	31
	Dedicated scanning tools	36
03	Summary/Conclusion	37

01

Intro

Explore the challenges of vulnerability management in cloud-native environments, and gain practical advice on how to improve your security posture. We cover topics such as identifying and prioritizing vulnerabilities, implementing vulnerability scanning and remediation strategies, and integrating vulnerability management into the DevOps workflow.

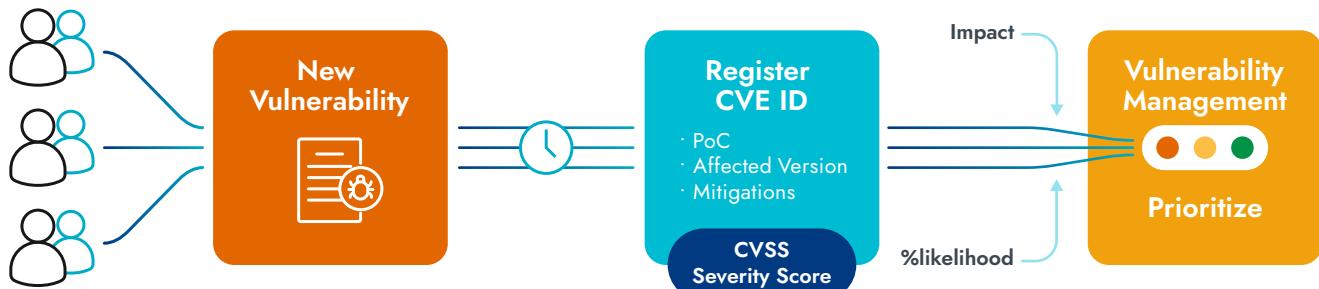
This book delves into the role of automation, the importance of continuous monitoring, and the need for collaboration between security and development teams. We will discuss the emerging technologies and best practices for managing vulnerabilities in Kubernetes and other container orchestration platforms.

This eBook serves as a comprehensive guide to evolving vulnerability management for cloud-native architectures. It will help security teams gain a deeper understanding of the unique challenges they face, and provide them with the knowledge and tools necessary to enhance their security posture and protect their organization's assets in the cloud.

Vulnerability Management Then and Now

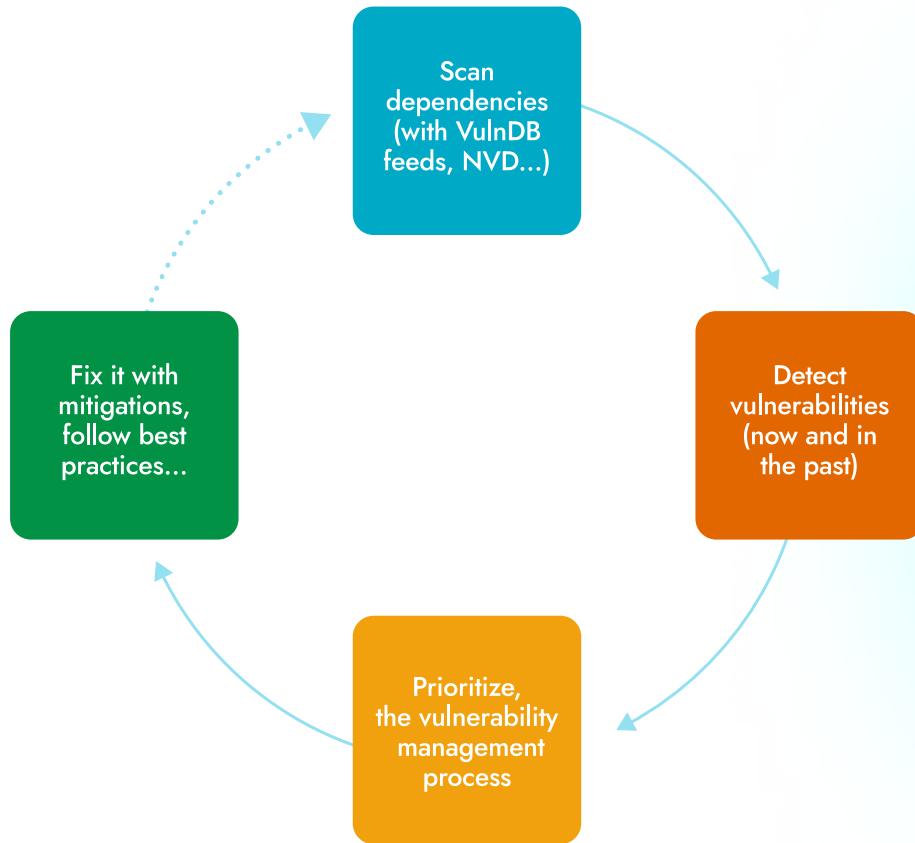
The approach to vulnerability management has evolved significantly over time. Previously, people relied on tools like Nessus to scan their data center networks and identify flaws in third-party software like operating systems, web browsers, and network device firmware. The quantity of flaws always exceeded the quantity of time available to fix them, and good approaches to prioritize the effort by actual risk to the organization did not emerge until fairly recently.

Today, in-house software development is becoming more prevalent every year. Businesses differentiate by creating better software, which gives rise to new challenges in vulnerability management. New stakeholders emerge as software development teams and DevSecOps teams enter the picture, and their incentives heavily prioritize speed of delivery. The number of flaws increases while the amount of time available to fix them decreases. Although conventional vulnerability management techniques are still applicable, this document primarily focuses on addressing the latest vulnerability management challenges associated with securing cloud-native applications delivered primarily via CI/CD pipelines.



The Framework

While the overarching framework for vulnerability management remains mostly unchanged, there are some notable differences in the modern software development landscape. The basic steps of vulnerability management still include asset identification, scanning, mitigation, remediation, and ongoing verification.



However, several new factors come into play when managing vulnerabilities in modern software development:

- **Speed of delivery is a business priority**

First, speed is a crucial consideration. In the fast-paced world of software development, organizations need to keep up with the latest technologies and push out new products quickly. This means that vulnerability management must not hinder the development process or slow down releases.

- **Developers are a key stakeholder in the vulnerability management program**

Second, vulnerability management in modern software development requires close collaboration with developers. Developers are the ones creating the software, so remediation often falls squarely on them and any mitigation requires their input. Collaboration between security and development teams is crucial to ensure vulnerabilities are found and addressed promptly. This relationship does not exist in traditional IT organizational structures.

- **Full-context risk prioritization is a non-negotiable requirement**

Third, vulnerability management must be informed by a thorough understanding of the context in which the software operates. This includes factors such as the application's purpose, the types of data it handles, and the potential impact of a vulnerability being exploited. Risk must be prioritized accordingly, and vulnerabilities that pose the greatest risk to the organization must be addressed first.

In conclusion, while the basic framework for vulnerability management has remained largely unchanged over time, the modern software development landscape requires a more nuanced approach. Organizations need to balance speed with security, involve developers in the process without hindering their agility, and prioritize risk based on the context in which the software operates.

Vulnerability Management for Software You Build

Cloud-native environments where applications are designed and executed at scale introduce us to some new operating models. Many of the underlying concepts stem from familiar ideas like “secure by design” development strategies and defense in depth. We’ll explore some of the unique nuances and challenges in this section.

Shift left

“Shift left” is a term used in the context of vulnerability management to refer to the practice of integrating security considerations to catch vulnerabilities earlier in the software development lifecycle (SDLC), rather than waiting until later in the process (usually at runtime) when they can be more difficult and costly to fix. This can help to improve the overall security posture of the organization and reduce the potential impact of security breaches.

Traditionally, vulnerability management has been focused on identifying and remediating vulnerabilities in deployed applications and systems, often as late as in production environments. However, this reactive approach can result in vulnerabilities that are more difficult to remediate and can create significant security risks, especially at scale.

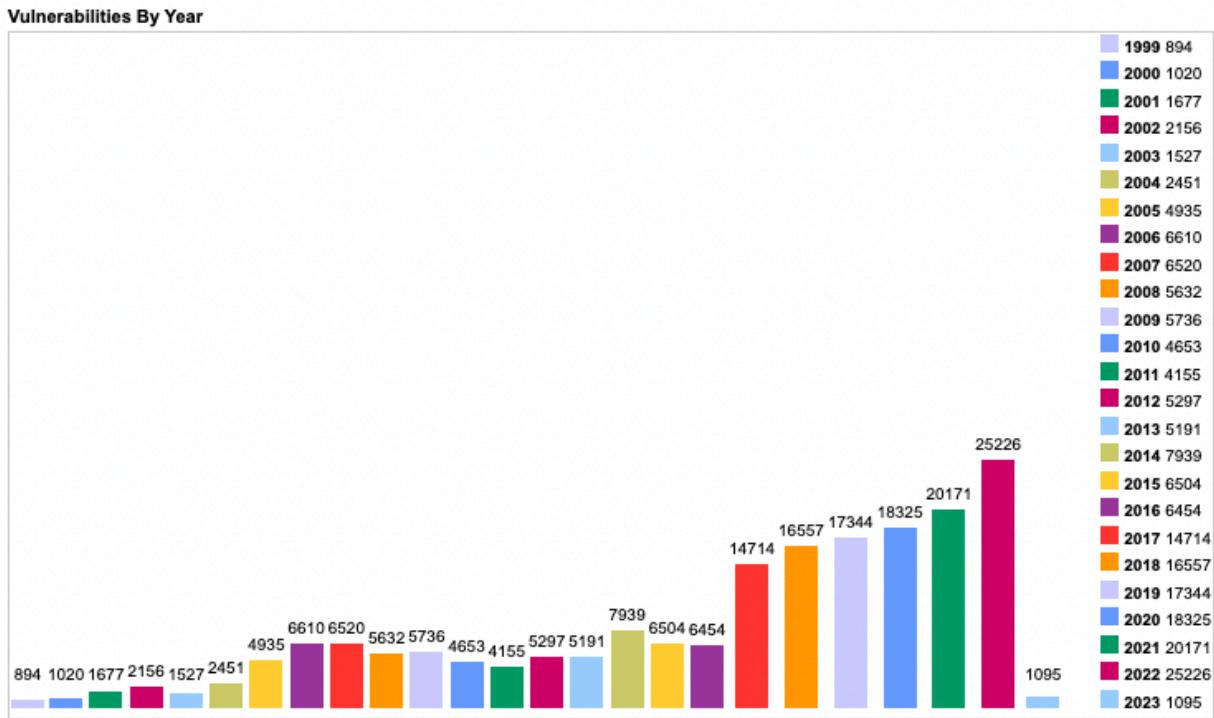
In rare cases, the remediation of a flaw can require significant rework of the application logic, placing a huge and avoidable burden on the engineering teams. By shifting left, organizations can address security issues earlier in the development process, which can help to reduce the number and severity of vulnerabilities that are introduced into production environments.

Full Lifecycle Vulnerability Management

Full Lifecycle Vulnerability Management is a crucial aspect of an organization’s security posture, and it involves managing vulnerabilities throughout the entire application lifecycle. This includes identifying, prioritizing, and remediating vulnerabilities in a timely and efficient manner.

Traditional vulnerability management programs tend to be scoped to servers, applications, network devices, and other data center assets. The program attempts to aggregate the findings from a single scanning tool (like Nessus) into a central dashboard for the security team to review. Remediation effort is handed off to the system owners via traditional workflow systems (like ServiceNow), who reside in various IT operations teams.

However, vulnerability management programs scoped to include cloud-native or any home-grown software will need to utilize additional tooling, such as Static Application Security Testing (SAST), Dynamic Application Security Testing (DAST), Software Composition Analysis (SCA), and developer-friendly workflow tools (like JIRA).



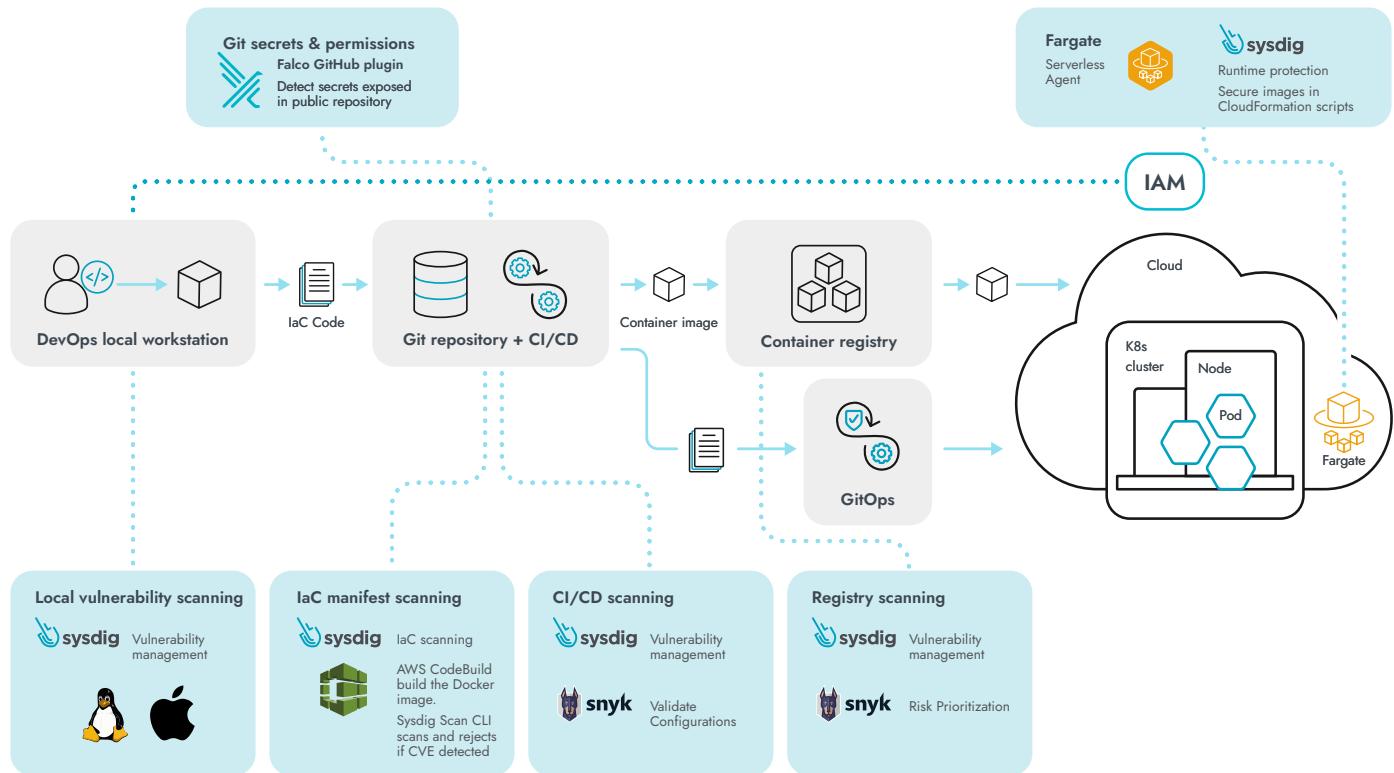
Risk prioritization of vulnerability assessment findings was always important because it's impossible to fix all of the flaws, especially as new ones are discovered every day. The challenge of knowing where to send those findings and how to ensure they are actionable is exacerbated by the complexity of modern architectures and the involvement of developers, who typically have very little security training.

Whenever possible, organizations should strive to prioritize vulnerabilities by risk within a scope of remediation accountability, rather than consolidating all findings into a single platform. Each application or asset owner should care about the highest risk vulnerabilities on their system. For example, software developers may not care about vulnerabilities on network switches unless they wrote the firmware.

Therefore, Full Lifecycle Vulnerability Management should not be a one-size-fits-all approach, but rather a customized and integrated approach for each software development organization. This can be accomplished by integrating vulnerability management tools into the DevSecOps pipeline, ensuring that vulnerabilities are identified early in the development process and remediated before they become larger problems.

In summary, Full Lifecycle Vulnerability Management should prioritize risk within a DevSecOps team's scope of control and seamlessly integrate with DevSecOps processes to ensure timely and efficient response.

Where are the vulnerabilities?



There are a variety of areas where vulnerabilities could be found, some of them are either nonexistent or not accessible in traditional Vulnerability Management programs. A modern VM program must scan every possible source of vulnerability introduction into the software and understand who is responsible for addressing these flaws.

Application code

The source code itself is where the vulnerability begins its life, and it is entirely in the software developer's scope of responsibility. Training, Security Champions programs, and "secure by design" methodology help ensure developers write more secure code. Application Security Testing (AST) tools like linters and SAST provide assessment capabilities for the code itself.

Application artifacts

These are vulnerabilities that exist within the artifact generated after the source code is packaged into its deployable form. For example, Jar files, which are commonly used to package Java applications, can introduce vulnerabilities if they contain outdated or vulnerable dependencies. Who owns the vulnerability assessment of this asset depends on the type of artifact in question and how the team is organized. Typically, developers are responsible for ensuring the artifact they ship to staging environments are defect-free to the extent specified by security policy. However, security teams should use automated tools to scan the application components for vulnerabilities through all stages of testing and staging.

OSS dependencies

Many applications rely on third-party open source software (OSS) components to function. These components can introduce vulnerabilities if they are not properly maintained or updated. Developers are responsible for keeping track of these dependencies and ensuring they are up-to-date and free from vulnerabilities. Security teams should scan for vulnerable OSS components throughout the CI/CD pipeline. Security and/or DevOps teams sometimes curate and maintain a registry of pre-approved OSS (and commercial) components for developers to use, disallowing the inclusion of any other packages.

Included libraries

Similar to OSS dependencies, included libraries can also introduce vulnerabilities if they are not properly maintained or updated. Developers are responsible for keeping track of the libraries they use and ensuring they are up-to-date and free from vulnerabilities. Security teams should scan for vulnerable libraries throughout the CI/CD pipeline.

Container images

Container images are a special type of application delivery vehicle that includes everything an application needs to function in a single package. They allow for easy and rapid portability between platforms and alleviate “works on my machine” dependency problems. Vulnerabilities can exist within multiple layers of container images. Developers are responsible for ensuring that the image layers containing the application artifacts and enabling architecture components are up-to-date and free from vulnerabilities before they are deployed.

However, IT operations, DevOps, or security teams may be responsible for maintaining the base image layers, which include the operating system. Security teams should also use dedicated container scanning tools to identify vulnerabilities in container images throughout the CI/CD pipeline. Most traditional vulnerability management tools do not provide any visibility into container images or running containers.

Hosts/Nodes

The underlying infrastructure that supports cloud-native applications, such as Linux servers or IaaS platforms, can also introduce vulnerabilities if they are not properly secured or maintained.

Improperly secured infrastructure supporting cloud-native applications, such as Linux servers or IaaS platforms, can introduce vulnerabilities. Breached containers usually have limited impact, but a breached host can be a huge security risk. Though container hosts are disposable, they remain a risk. Immutable OSs, like CoreOS, RancherOS, RHEL Atomic Host, Ubuntu Core, etc., are image-based, safer, and more focused on running containers. Operations teams are typically responsible for securing underlying infrastructure. Special OSs for running containers may not be scannable by legacy tools.

Orchestration Systems (Kubernetes)

In traditional monolithic architectures, each application typically runs on its own dedicated infrastructure with its own set of permissions. This means that if a vulnerability is discovered in one application, it is less likely to affect other applications running on the same infrastructure. However, in a Kubernetes environment, multiple applications may run on the same infrastructure, and a compromised plugin can potentially affect multiple applications at once.

Kubectl plugins are a gateway for vulnerabilities

Kubectl plugins are external binaries that can be invoked through the kubectl command-line tool to extend its functionality. When a user runs a kubectl plugin, it runs with the same permissions as the kubectl command itself. This means that if the kubectl command is run with root privileges, the plugin will also have root privileges.

This can be a unique problem for the Kubernetes orchestration platform because Kubernetes is designed to run multiple containers on a single host, each with its own set of permissions. If a kubectl plugin is compromised, it can potentially access or modify any resource within the cluster, including other containers or nodes. Are these plugins actively maintained? If there's a CVE in one of these plugins, how do we handle these unpatchable vulnerabilities?

To address this challenge, it is important to carefully manage the permissions granted to kubectl and its plugins. This may involve limiting the use of plugins to trusted sources or auditing the code of plugins before they are used in production environments. Additionally, vulnerability scanning in Kubernetes environments should be a continuous process, with regular scans to identify and mitigate any vulnerabilities that are discovered.

Unpatchable vulnerabilities

The issue with unpatchable vulnerabilities in Kubernetes, such as [CVE-2020-8554](#), is that they represent a potential security risk that cannot be fully eliminated through traditional patching or remediation methods. These vulnerabilities are often the result of fundamental design flaws or limitations in the technology itself, which cannot be easily or quickly addressed.

In cases where there are known, unpatchable vulnerabilities in Kubernetes, there are a few steps that organizations can take to protect themselves. One option is to use a tool like Open Policy Agent (OPA) to enforce strict policies around the use of Kubernetes resources and configurations. By setting up policies that restrict certain actions or configurations that could lead to exploitation of the vulnerability, organizations can reduce the risk of an attack.

Alternatively, you can monitor the environment for signs of exploitation or compromise using a tool like Falco. This can involve setting up alerts and monitoring systems to detect any unusual activity or changes in the Kubernetes cluster, as well as conducting regular security audits to identify potential vulnerabilities.

Unfortunately, the best approach will depend on the specific vulnerability and the risk it poses to the organization. In some cases, it may be possible to mitigate the risk through policy enforcement and monitoring, while in other cases, it may be necessary to take more drastic measures, such as limiting or disabling certain features or functions within Kubernetes. It is important to work closely with security experts and the Kubernetes community to stay up-to-date on the latest vulnerabilities and mitigation strategies.

Shared responsibility model of IaaS offerings

The shared responsibility model in the cloud refers to the fact that cloud providers are responsible for the security of their underlying infrastructure, while customers are responsible for the security of their own applications and data that resides on the cloud infrastructure.

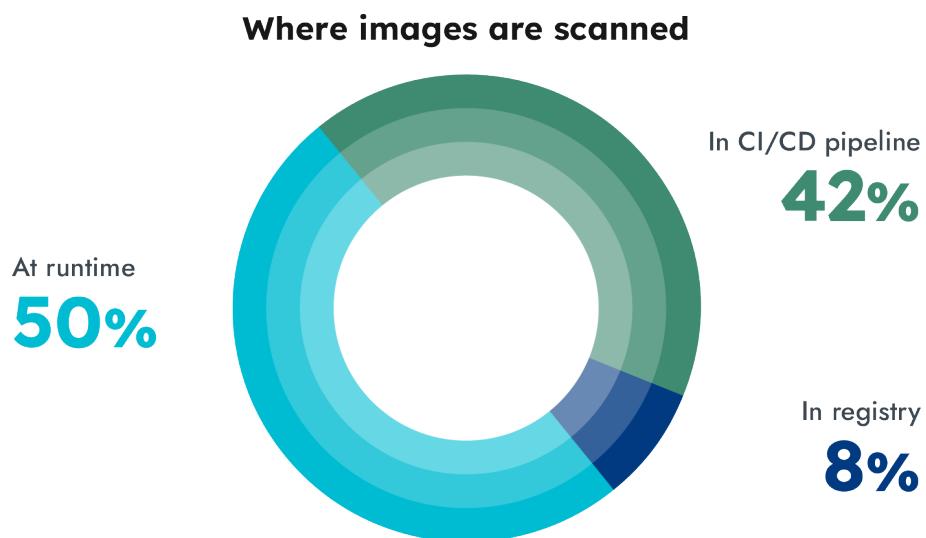
While this model can be beneficial in many ways, it can also lead to some challenges when it comes to vulnerability detection and mitigation. One example is the case of managed Kubernetes services, like GKE, EKS, and AKS. While these services provide a convenient and scalable way to deploy and manage Kubernetes clusters, they also limit the control that customers have over the underlying infrastructure. This can make it challenging to detect and mitigate vulnerabilities that may exist in the infrastructure or in the Kubernetes cluster itself.

One approach to addressing this challenge is to use third-party vulnerability scanning tools that are specifically designed to scan Kubernetes clusters. These tools can help to identify vulnerabilities in both the infrastructure and the applications running on the cluster, and can provide recommendations for how to mitigate those vulnerabilities.

We recommend working closely with the cloud provider to ensure that they are meeting their responsibilities for securing the underlying infrastructure. For example, you can ask for regular security audits or penetration testing to be performed on the infrastructure to ensure that it is secure. In reality, the cloud provider can do a “best effort” at updating vulnerable components, but likely cannot make radical changes for all customers based on individual feature requests.

Third-party architecture components (Middleware)

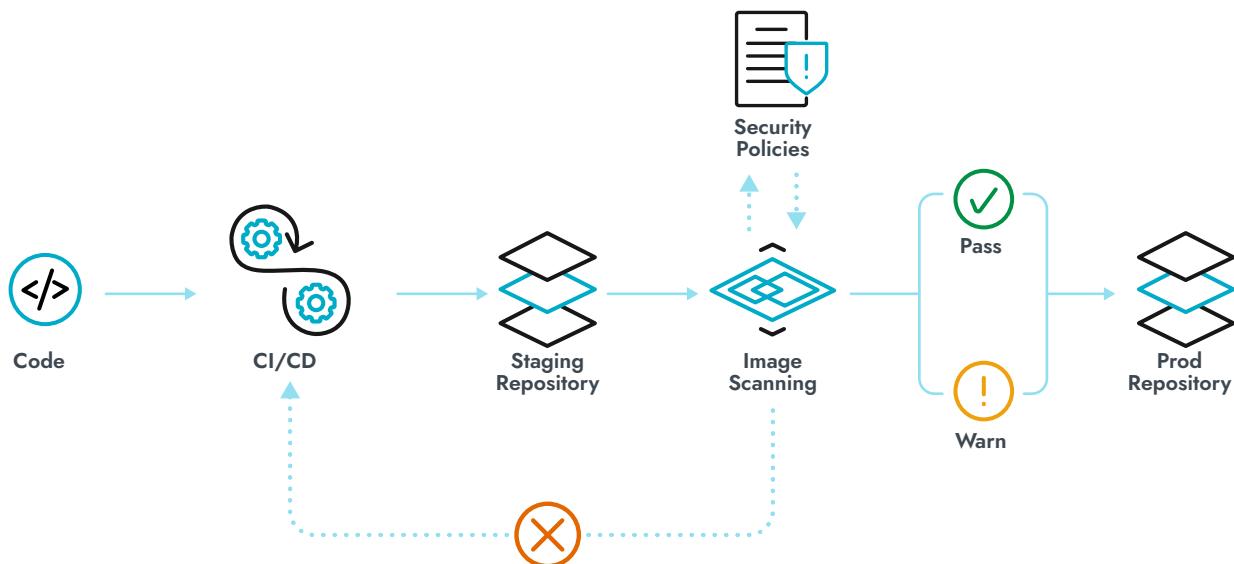
Cloud-native applications often rely on third-party architecture components, whether OSS or commercial. These are essentially modern middleware and include things like load balancers, databases, and much more. These components can also introduce vulnerabilities if they are not properly maintained or updated. Developers are responsible for keeping track of these components and ensuring they are up-to-date and free from vulnerabilities. Because they are not built in-house, these pieces of software often bypass the CI/CD security controls (and often the pipeline altogether) and may not be assessed until runtime. According to the [2023 Sysdig Security Usage report](#), only 42% of vulnerability scans are performed at the CI/CD pipeline phase. Ideally, security teams will scan for vulnerable architecture components throughout the CI/CD pipeline.



To recap, identifying and fixing vulnerabilities requires a collaborative effort between different teams, including developers, security professionals, and operations teams. By staying vigilant and regularly scanning for vulnerabilities, organizations can help ensure the consistent security of their cloud-native applications.

When are the vulnerabilities discovered?

Based on the above information, when is it appropriate to scan for vulnerabilities in cloud-native environments? The short answer is through every stage of the CI/CD pipeline and then constantly in production to account for any disclosures that occur after deployment.



On developer machine before merge

One best practice is to scan for vulnerabilities on the developer's machine before the code is merged with the main branch. This can help catch vulnerabilities early in the development process, when they are typically easier and cheaper to fix. Developers can use tools such as dependency checkers and static analysis tools to scan their code for vulnerabilities.

In CI/CD pipeline

Another best practice is to incorporate vulnerability scanning into multiple stages of the continuous integration and continuous deployment (CI/CD) pipeline. This ensures that code is scanned for vulnerabilities each time it is built and deployed. Scanning should occur at multiple points in the pipeline, including during code builds and before the application is deployed. This can introduce a high degree of automation to vulnerability discovery, which is good for supporting delivery speed. However, the tailoring of pass/fail gates can be difficult and must be carefully considered in the context of each organization's specific risk appetite.

Registries

Cloud-native applications are often deployed using container images, which are stored in registries. Non-containerized artifacts also live in registries. Most organizations have many registries and many different types of registries. Before deploying an artifact, it's important to scan it for vulnerabilities. Registry scans can occur multiple times in the Lifecycle, but it's absolutely critical to ensure the production registry is clean.

Runtime

While it's best to catch vulnerabilities as early as possible, it's also important to scan for vulnerabilities at runtime. This can help identify vulnerabilities that were not caught during earlier stages of the development process, those that may have been introduced during runtime, or those disclosed after the last scan occurred. Runtime scanning should be done regularly to ensure that any vulnerabilities are identified and addressed promptly. It's worth noting that "runtime" can mean staging or production, and ideally, the teams have full scanning capability in both types of runtime environments.

We cannot stress enough the importance of incorporating vulnerability scanning into the various stages of the development process. This will ultimately help to ensure the security of your cloud-native environments. By catching vulnerabilities early and regularly scanning for them, organizations can help minimize the risk of cyber attacks and protect their sensitive data.

Prioritization of Vulnerability Assessment Findings

Arguably one of the hardest aspects of vulnerability management is prioritization, as not all vulnerabilities have the same level of risk, and resources must be allocated efficiently to address the most critical ones first. This guide provides a comprehensive framework for categorizing vulnerabilities based on their potential impact on your organization's assets, services, and operations. By following these guidelines, you can prioritize your vulnerability management efforts and ensure that your teams are focused on the most critical issues.

Risk contextualization is an approach taking into account the specific context and environment in which vulnerabilities exist. It involves analyzing the severity of the vulnerability, the assets at risk, the potential attackers and their motivations, the available mitigation options, and more.

We aim to highlight the importance of risk contextualization in prioritizing targeted controls to mitigate workload and cloud-based threats. We use a two-part framework to think about vulnerability risk:

- **Vulnerability Threat Context** comes from the threat intelligence surrounding the vulnerability itself, regardless of which assets, if any, it is affecting. This includes information like whether an exploit is publicly available or likely to become available soon, whether the exploit is easy to execute, and whether the exploitation of this flaw is popular among attackers today. Threat context data usually comes from third-party sources, and could be in machine-readable form or in the form of a report. Many vulnerability assessment tools integrate threat intelligence sources to help prioritize the scan results within their interface. Additional threat context can be manually gathered by security teams from public and commercial sources like [Exploit Database](#), [Metasploit Framework](#), and others.
- **Affected Asset Business Context** considers the importance of the assets within your organization that can be affected by any given vulnerability and the potential impact to your business should that asset be abused. Additionally, this includes the details of your environment, such as whether the asset is accessible from the internet or what mitigating controls may be in place to protect it.

In summary, risk contextualization enables organizations to prioritize vulnerabilities based on their potential impact on the business and the likelihood of successful exploitation in the specific context of their environment. This approach helps in optimizing resource allocation, reducing risk exposure, and enhancing the overall security posture of the organization.

There are dedicated frameworks such as the [Stakeholder-Specific Vulnerability Categorization \(SSVC\)](#), a vulnerability management methodology that assesses vulnerabilities and prioritizes remediation efforts based on exploitation status, impacts to safety, and prevalence of the affected product in a singular system.

There are also frameworks for calculating a risk score like EPSS, which can help your team quantify your own risk assessments. However, these are often cumbersome to use and results can be inconsistent. Ideally, commercial vulnerability assessment tools have risk-based prioritization built in and tied to an actionable remediation workflow.

Runtime risk intelligence for cloud-native applications

The methodology above applies to any vulnerability management program, but when handling home-grown software, we have some additional considerations. Building software packages with only necessary dependencies is important for reducing complexity and minimizing the risk of vulnerabilities. However, even with careful selection of dependencies, it's possible that some unnecessary ones may slip through the cracks. Operating System (OS) packages, in particular, include a lot of unnecessary components, which may contain vulnerabilities, but are never used by your application. In fact, the 2023 Sysdig Cloud-Native Security & Usage Report stated that OS packages contained 37% of the vulnerabilities. This can lead to bloat in the package and make it difficult to identify which dependencies are actually being used.

A recent innovation in cloud-native risk prioritization is [Risk Spotlight](#), which provides visibility into which dependencies are actually being used by your application in runtime and prioritize them based on their real risk. This allows your development teams to focus their efforts on addressing the most critical vulnerabilities and reducing the impact of cyber attacks on your business. By eliminating unnecessary dependencies and prioritizing those that are actually being used, you can streamline your software package and improve the security and stability of your applications.

Remediation

At some point, you will need to either fix or mitigate vulnerabilities in an organization's systems and applications. Effective remediation requires identifying and prioritizing vulnerabilities based on their severity and impact, as discussed in the previous section. We already mentioned that the teams responsible for performing the remediation can vary depending on which environments, applications, or even individual components the vulnerability affects.

Asset inventory

The first step to enabling a smooth workflow here is maintaining an accurate asset inventory with an owner assigned to every entry. Because cloud-native applications tend to be highly distributed and complex, keeping track of them isn't easy.

For the context of ephemeral workloads, tagging can help organizations keep track of their assets in real time. With ephemeral workloads, assets are constantly created and destroyed, which can make it difficult to maintain an accurate inventory. By utilizing tags, you can quickly identify which assets are currently in use, and which ones have been decommissioned, as well as who the application owners are and what constraints or requirements may be associated with the workload.

Tagging can also be used to enforce compliance checks. You could assign specific tags to assets that must comply with strict security requirements, such as access controls or encryption. In the case of vulnerable images, we can state that these do not comply with those rigorous compliance standards, and therefore cannot be used in FedRAMP environments.

Remediation responsibility

Who is actually capable of performing the remediation has never been simple to answer, but the nature of cloud-native applications exacerbates the problem further. For example, different layers of a container image can be the responsibility of different people or teams. An IT Ops or DevOps team might own the operating system layers, but a development team might own the application layers.

Similarly, fully custom built components of an application may be owned by developers while third-party components and middleware might be owned by another team. It's critical that the security function that performs the scanning hands-off remediation responsibility to the right people and in a rigorously risk-prioritized way, with the understanding that Dev and Ops teams do not usually have the security expertise to properly prioritize remediation work.

Metrics and incentives

A key aspect of secure vulnerability management is ensuring that remediation efforts are timely and effective in addressing vulnerabilities that could be exploited by attackers. However, different technical teams involved in this process may have competing metrics and incentives. Developers are rewarded for shipping a lot of features as quickly as possible because that ultimately creates value for the business.

Security, on the other hand, must protect the business from harm, so each flaw that slips into production is an additional source of risk. As such, modern organizations must leverage automation, risk-based prioritization, and new tools like admission control to optimize the release of safe software quickly. Ultimately, a culture shift is required, and security and engineering leaders alike must incentivize their teams to work together to avoid the adversarial relationships that plagued traditional vulnerability management.

Mitigation

Vulnerability mitigation is an important step in the vulnerability management process, where identified vulnerabilities are temporarily addressed to minimize their impact while awaiting a more permanent solution. In some cases, remediation of the vulnerabilities may require significant disruption to the system or process, and therefore, mitigation is employed as a temporary holdover until the remediation can be properly executed.

In an ideal world, this would mean **fixing it**. Updating the package or library to a version that contains the fix for the aforementioned vulnerability in the source code (patch) and then trigger a new deployment for that particular piece of software: build the artifact, perform the tests required, etc., and finally deploy it into the production environment.

A significant hurdle to overcome with respect to remediation is to be able to quickly patch every single asset or dependency that is impacted or potentially exploitable, and these processes also need to be able to scale. It's not trivial to patch old versions that could impact new features or poor performance; do it in a massive way or have other implications. The problem is exacerbated with transitive dependencies. That is, your code or system likely relies on many other codebases or systems, and dependency chains become quite nested in practice.

Being able to include as much context as possible for the developers to fix the vulnerability is important to avoid wasting time on figuring out who is responsible or even how to fix it.

Unfortunately, it is not always possible to fix a vulnerability. Maybe the fix is not available yet, requiring more time to test that it doesn't break anything else, or any other million reasons.

There are a few approaches if that's the case:

- **Downtime.** If the potential breach is big enough, maybe the affected application shouldn't be running. When the [Log4j vulnerability](#) was found, [Quebec shut down 3,992 websites as a 'preventative measure.'](#)
 - If the vulnerability is focused into a specific feature, maybe just that particular feature can be disabled. A proper way to do that would be using [feature flags](#).
- **Hardening.** Try to prevent the exploitation of the vulnerability by making some changes to the application or the system if possible.
 - If the vulnerability can be network exploitable, maybe enforcing the [security groups](#) or the [Kubernetes network policies](#). Here is a network policy to block all the egress traffic for a specific namespace:

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
Metadata:
  name: default-deny-egress
  namespace: default
Spec:
  podSelector: {}
  policyTypes:
  - Egress
```

- **Monitoring.** Monitor the system for suspicious activity and respond when necessary.

This last step overlaps with a whole security category called "[Threat detection & response](#)." By using a Runtime detection engine tool like [Falco](#), you can detect attacks that occur in runtime.

Let's assume that the attacker exploits a vulnerability and wants to open a reverse shell on a pod. In this case, the Falco runtime policies in place will detect the malicious behavior and raise a security alert.

Validation

A clear process associated with remediation is confirming that a vulnerability has been successfully remediated. To do this, you will need to verify that the remediation efforts were successful to ensure that the vulnerability is no longer present, and that the organization's systems and applications are secure.

Verification can be done through various means, such as manual testing, automated scanning, or third-party assessments. Most tools' reporting capabilities allow you to run a "before-and-after" report to verify a highly-critical CVE is no longer present in your environment.

Risk acceptance

Sometimes, it is not possible to fix a vulnerability because it may not have a fix already available, will take some time because of organizational processes, or simply because it is a false positive.

In those scenarios, it can be helpful making a conscious decision to filter out ("accept") those vulnerabilities so they don't affect the report results, but it is important to have a clear understanding of the consequences and to put in place some methods (such as reminders) to review if the filter can be removed.

Reporting

Reporting the results of vulnerability assessment is a key part of this process, and it will take on a different form depending on who the target audience is. All scanning tools should be able to create some kind of report, but their content and structure is often inconsistent.

A key consideration is that reporting to senior leadership should be far less granular than reporting of technical teams. Executive reports must include high level program metrics, preferably showing trends over time. Technical reports should be structured in a way that is most useful to the team receiving the information. Effective reporting requires clear and concise communication that highlights any problems within the vulnerability program, not just a list of identified CVEs.

Nuances of Cloud-Native Architecture

Cloud-native architectures bring several additional nuances when it comes to vulnerability scanning. These unique considerations stem from the nature of modern architectures and the increasing degree of automation that IT requires today.

- **High degree of environment complexity**

Cloud-native architectures rely on microservices, which are small, independent components that perform a specific function. This can make it more difficult to scan for vulnerabilities because the components may be spread across multiple hosts or data centers and may have different security requirements.

- **Ephemeral nature of infrastructure**

Cloud-native architectures are designed to be ephemeral, meaning that the clusters, hosts, pods, and containers are intended to be created and destroyed regularly and frequently. Infrastructure components are provisioned and deprovisioned dynamically in response to system or customer demand. This can pose challenges for vulnerability scanning because the components may not exist long enough to be scanned. The IP and Container ID will change, and therefore cannot be tracked easily.

- **Immutable components**

Cloud-native infrastructure and application components should be immutable. This means that changes to the infrastructure should be made by creating new instances rather than modifying existing ones. This approach can make it more difficult to scan for vulnerabilities because the infrastructure components are constantly being replaced.

- **Massive scalability**

Cloud-native architectures are designed to be elastic and scalable, which means that they can easily be expanded or contracted as needed. This can pose challenges for vulnerability scanning because the infrastructure components may be spread across multiple hosts or data centers.

- **Workflow automation**

Cloud-native architectures are designed for automation to manage infrastructure components quickly and scalably. This automation can make it more difficult to perform vulnerability scanning because the scanning process may need to be integrated into the automation tools.

This section explores some additional details that can either impede or bolster your vulnerability management program. Not every program can address every nuance from the beginning, but as your organization increases its DevSecOps maturity, these elements will certainly become relevant.

Integrating Into the Development Lifecycle

Scanning for vulnerabilities is a [best practice and a must-have](#) step in your application lifecycle to prevent security attacks. It is also important where this step is performed, but why?

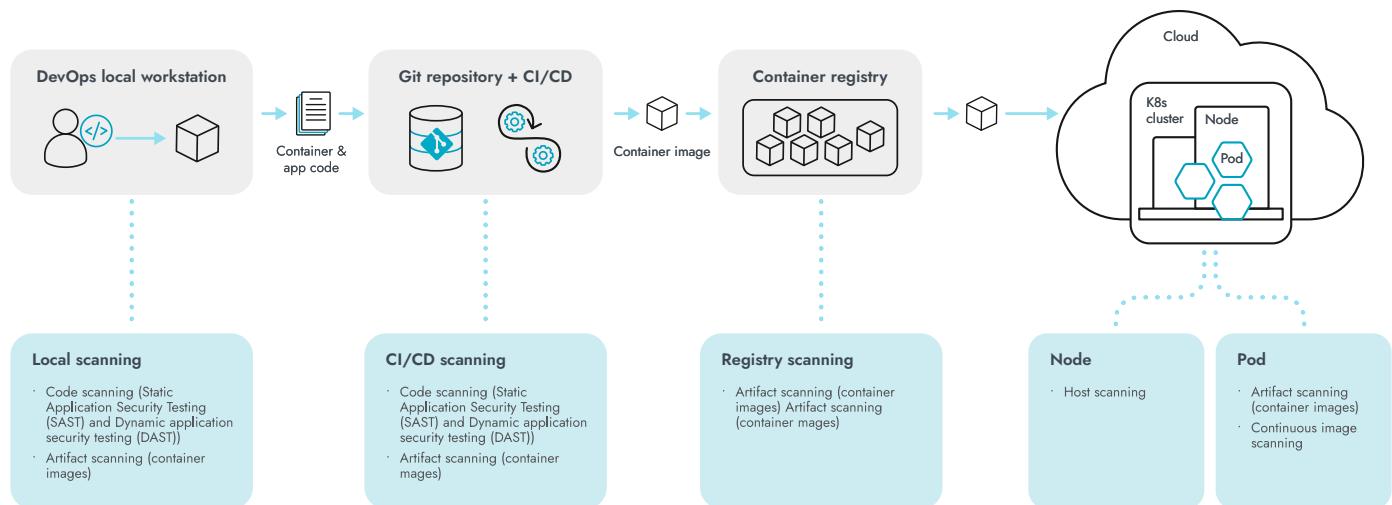
Applications' lifecycle involves a number of steps, from the developer workstation creating fine art in the shape of lines of code to the final production environment where the customers use a web application, mobile application, or anything else. Vulnerabilities can be introduced in any of those steps, so it is highly recommended to put some barriers to prevent them from ruining your environment.

The "[defense in depth](#)" concept recommends performing automatic vulnerability scanning on different steps of the application lifecycle (sometimes even overlapping them). This will reduce the number of vulnerabilities introduced in your production environment.

In today's ever-evolving threat landscape, it's becoming increasingly clear that vulnerability scanning before production is not sufficient. This is because new vulnerabilities are constantly being disclosed, and unexpected behavior can occur during runtime. As such, vulnerability scanning processes should adopt a "complete or comprehensive vulnerability assessment" approach, which emphasizes the importance of thoroughly assessing vulnerabilities without any preconceived trust assumptions.

What if the CI/CD was bypassed and someone pushed the container image directly? What about those images that were scanned weeks ago? Were new vulnerabilities discovered since that last scan?

It is also important to note that fixing vulnerabilities earlier in the software development process is easier and that every step of the process makes things more complicated. That's why "shifting left" security improves the overall security posture of the organization and reduces the potential impact of security breaches.



CI/CD

Let's assume you already fixed all those vulnerabilities by updating the libraries dependencies and the PR has been submitted. The next step in the build chain is usually running a [CI/CD pipeline](#) to build the application, build the container image, run some tests... and check for vulnerabilities again. Wait, what? Why again?

- Who can guarantee the vulnerability scan has been done religiously by all the developers locally in their workstations before submitting the pull request?
- What if the developer performed the scan a couple of days ago and a new vulnerability has been found?

CI/CD pipelines are basically made of different steps.

A very basic example can be something like:

- Checkout the code
- Run some linting to make sure the build won't fail
- Build the artifact
- Perform some unit tests
- Deploy the artifact

But they can be as complex as needed, it depends on the requirements of the application or the environment itself. The benefit of this flexibility is that there can be as many steps as needed – even multiple steps can happen at the same time if needed. Adding as many security checks as possible in a CI/CD pipeline is a good idea.

Let's see a more complex example. This time, the application is packaged as a container image:

- Checkout the code
- Run some linting to make sure the build won't fail
- Check for vulnerabilities on the dependencies
- Check for misconfigurations or secrets
- Static code scan
- Build the application and the container image
- Check for vulnerabilities on the container image
- Perform some unit tests
- Deploy the artifact

Why check for vulnerabilities on the container image **again**? The answer is simple. What if the image used as a base already has vulnerabilities?

Ultimately, the goal is to deploy workloads to the production environment that are as clean as possible and within the company's security and compliance standards. CI/CD enables a lot of automation of the security assessment and remediation of any software your organization builds at speeds that were previously unimaginable.

Binaries, packages, and language intricacies

Package managers are the common approach to install software in container images or host operating systems. Tools like apt or dnf make it really easy to install, update, and manage the software. Most of those tools use a database to store the metadata about the packages, such as when it was installed, the dependencies, the versions, the files included in a package, etc.

Using the package version, it is trivial to check against the vulnerability databases to see if it is vulnerable or not. In the following example, we check the curl version in a Debian system and manually check if the version is affected by any vulnerability in the Debian vulnerability database (<https://security-tracker.debian.org/tracker/source-package/curl>):

Information on source package curl

curl in the Package Tracking System | curl in the Bug Tracking System | curl source code | curl in the testing migration checker

Available versions	
Release	Version
buster	7.64.0-4+deb10u2
buster (security)	7.64.0-4+deb10u5
bullseye	7.74.0-1.3+deb11u3
bullseye (security)	7.74.0-1.3+deb11u7
bookworm	7.88.1-1
sid	7.88.1-5

Open issues					
Bug	buster	bullseye	bookworm	sid	Description
CVE-2023-23915	fixed	vulnerable (no DSA, ignored)	fixed	fixed	A cleartext transmission of sensitive information vulnerability exists ...
CVE-2023-23914	fixed	vulnerable (no DSA, ignored)	fixed	fixed	A cleartext transmission of sensitive information vulnerability exists ...
CVE-2022-43551	fixed	vulnerable (no DSA, ignored)	fixed	fixed	A vulnerability exists in curl <7.87.0 HSTS check that could be bypassed ...
CVE-2022-42916	fixed	vulnerable (no DSA, ignored)	fixed	fixed	In curl before 7.86.0, the HSTS check could be bypassed to trick it in ...

Open unimportant issues					
Bug	buster	bullseye	bookworm	sid	Description
CVE-2021-22923	vulnerable	vulnerable	fixed	fixed	When curl is instructed to get content using the metalink feature, and ...
CVE-2021-22922	vulnerable	vulnerable	fixed	fixed	When curl is instructed to download content using the metalink feature ...

Resolved issues	
Bug	Description
CVE-2023-23916	An allocation of resources without limits or throttling vulnerability ...

Source: [Debian.org](https://security-tracker.debian.org/tracker/source-package/curl)

But what about binaries that are not installed using package managers? Golang binaries are the classic example of applications that are just a single binary copied into a container image. Something like this:

```
## Build
FROM golang:1.16-buster AS build
WORKDIR /app
COPY go.mod ./
COPY go.sum ./
RUN go mod download
COPY *.go ./
RUN go build -o /helloworld
```

```
## Deploy
FROM gcr.io/distroless/base-debian10
WORKDIR /
COPY --from=build /helloworld /helloworld
USER nonroot:nonroot
ENTRYPOINT ["/helloworld"]
```

There are no package managers involved, instead a helloworld binary has been compiled from source code and then copied into the final container image. How can we check if it included any vulnerable dependency?

In this particular example of using go, the `go build` command embeds some information about the dependencies (debug information) into the binary itself by default (it can be disabled), so they can be extracted (using <https://pkg.go.dev/debug/buildinfo>, for example) and then perform the matching against a vulnerability database. For golang, that would be the [Go Vulnerability Database](#).

Some other programming languages don't include that information and it requires to know the dependencies they are using implicitly (the package.json file for nodeJS for example).

The best approach to solve this problem would be that every piece of software included in the container image includes their own Software Bill of Materials (or SBOM), so the vulnerability scanning tools would easily extract them and compare them against the vulnerability databases.

Unfortunately, SBOM is frequently unavailable or incomplete. In this case, you must rely on dependency mapping tools (such as Checkov) and ensure that your teams understand the particular behaviors of the languages and package managers that they use.

Software Bill Of Materials

Software Bill of Materials (SBOMs) are useful for vulnerability management because they provide an inventory of software components used in an organization's systems and applications, enabling the identification of vulnerabilities and prioritization of remediation efforts. SBOMs include a detailed list of software components, including version numbers and dependencies, which can be integrated with vulnerability scanning tools to automatically identify vulnerable software components.

By having a complete inventory of software components, security teams can more accurately assess the risk of software vulnerabilities and identify potential vulnerabilities before they are exploited. Ideally, every piece of software should generate and publish its own SBOMs, but tools like `syft` can generate them from file systems or container images.

However, creating and maintaining an accurate SBOM can be challenging in complex software environments with many dependencies. Some components may be buried deep within the software stack or may be included in multiple places, making it difficult to identify and track them all. Additionally, software components may change frequently, and it can be challenging to keep the SBOM up to date.

While SBOM is a useful concept in theory, it may not be practical for all organizations running cloud-native workloads. Instead, dependency mapping tools can help organizations identify and track dependencies across the entire software stack, including open source libraries, and manage and mitigate vulnerabilities more effectively.

Update in flight or destroy and redeploy

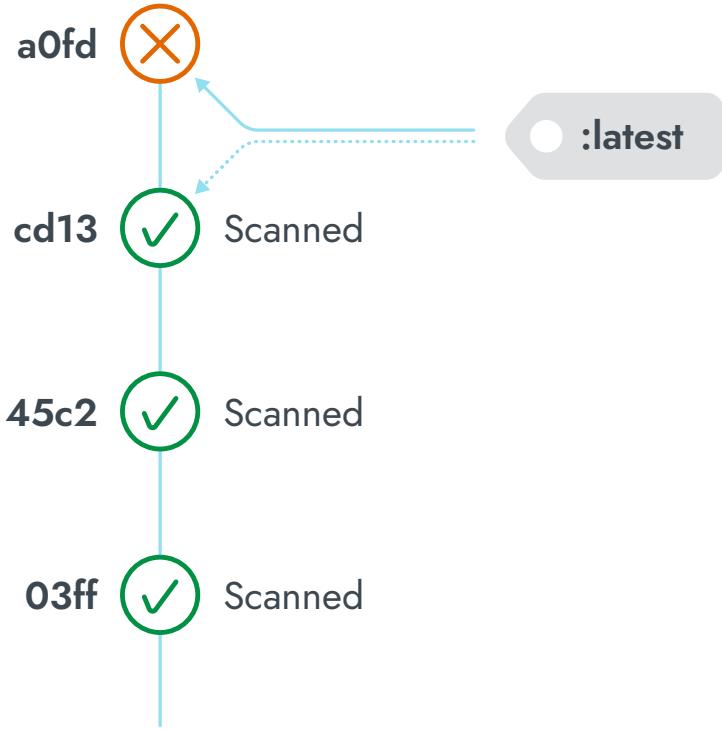
Traditional IT management likes to touch running systems to make necessary changes and updates in runtime. This goes against the immutable philosophy of cloud, but old habits die hard and not all applications in the cloud are fully cloud-native on the day they land there. For example, many cloud migrations include an application refactoring effort that can take a long time. In the early stages, the application may have a monolithic core with microservices around it as components are broken off into a more distributed architecture. In this case, the core may be managed in a "legacy" way while the newer microservices are treated as immutable.

Containers, in particular, were intended by design to be strictly immutable, which is why images are rebuilt and redeployed any time a change needs to be made. But is this always the right approach?

Using the latest package versions is always a good idea because it should contain less vulnerable software. The main concern when updating the container image packages is if by updating the packages/dependencies, the container image behavior or the application changes (breaks, or behaves differently). To avoid this, having a proper lifecycle process to ensure it is properly tested before the container image is running on the production environment is a must.

Pinning image versions

Sometimes, the image you scan is not the same one you deploy in your Kubernetes cluster. This can happen when you [use mutable tags](#), like “latest” or “staging.” Such tags are constantly updated with newer versions, making it hard to know if the latest scan results are still valid.



Source: [Sysdig Blog](#)

Using mutable tags can cause containers with different versions to be deployed from the same image. Beyond the security concerns from the scan results, this can cause [problems that are difficult to debug](#). Imagine the following scenario:

- An application was already deployed using the “:latest” tag two weeks ago and it is running properly.
- Suddenly, a traffic peak requires to deploy another instance of that same application on another node that didn’t have the container image cached, so it will pull the “:latest” one... which depending on when the latest build process has been performed, will include different package versions!

For example, instead of using `ubuntu:focal`, you should enforce the use of immutable tags like `ubuntu:focal-20200423` when possible.

Keep in mind that version tags (for some images) tend to be updated with minor, non-breaking changes. So, although it looks a bit verbose, the only option to ensure repeatability is using the actual image ID:

```
> ubuntu:@  
sha256:d5a6519d9f048100123c568eb83f7ef5bfkad69b01424f420f17c932b00dea76
```

That image ID is unique (it is calculated using the sha-256 algorithm) and identifies the container image in a way that cannot be tempered.

```
> podman images --digests  
REPOSITORY          TAG      DIGEST  
IMAGE ID            CREATED   SIZE  
docker.io/library/busybox  latest  
sha256:7b3ccabffc97de872a30dfd234fd972a66d247c8fc69b0550f276481852627c  
abaa813f94fd  2 months ago  3.96 MB  
docker.io/kindest/node  <none>  
sha256:f52781bc0d7a19fb6c405c2af83abfeb311f130707a0e219175677e366cc45d1  
476b7007f4f5  4 months ago  828 MB
```

Container build caching

A fully immutable approach would necessitate rebuilding the entire image from start to finish every time something changes. However, some images are very large and have quite a complex build process. In this case, image caching can be used to avoid rebuilding all of the enabling layers that are not affected by the change. Caching increases how quickly you can update, but it can be dangerous.

Consider the following Dockerfile:

```
FROM ubuntu:focal  
RUN apt-get update && \  
    apt-get upgrade -y && \  
WORKDIR /  
COPY ./helloworld /helloworld  
USER nonroot:nonroot  
ENTRYPOINT [ "/helloworld" ]
```

The first problem is that if you don't re-pull the base image (`ubuntu:focal`), it will be reused on every build. Fortunately the docker build command includes the `--pull` flag to always attempt to pull a newer version of the image.

The next problem is that the first build will download a few Ubuntu packages instructed by the `apt-get` command. However, the next runs of the build process will use the cached layers, including the one running the `apt` commands:

```
$ podman build .
STEP 1/6: FROM ubuntu:focal
STEP 2/6: RUN apt-get update &&      apt-get upgrade -y
--> Using cache
25ccbef2ada488518c7caca96315207d4f09c9657940ef412db1acdd8f3573a4
--> Pushing cache []:
0999c867cd96ff3ef67eac04ac945b942e97efada838892a4aa3760d1fda04eb
--> 25ccbef2ada
STEP 3/6: WORKDIR /
--> Using cache
1338e125b9de15f8dbf38b3a5d27cc34729d2a0c08a469dd3fa29f07d91207f7
--> Pushing cache []:
0ff44936c0286aff66b5d38d25faeeac8cb7931c6c89bb67c8b539c741710373
--> 1338e125b9d
STEP 4/6: COPY ./helloworld /helloworld
--> Pushing cache []:
ec902a775d3c721ce12529b83febaa87a3889ce2e6cbb5ec9253183027e6ba5c
--> dc512d2e412
STEP 5/6: USER nonroot:nonroot
--> Pushing cache []:
f5e2a671d74cb39d790cef908d20361fafae86720d608a3cb6dd2d1c48f5c18d
--> edb751c934a
STEP 6/6: ENTRYPOINT ["/helloworld"]
COMMIT
--> Pushing cache []:
b015981c164c99e5a58ed6c4cafaac34432016530d2b984d4c324cf93f2b358d
--> dc512d2e412
af109fa0ff21b01e575bbfe034ebcd74068697dafbeabd6323945451543b0102
```

Unless the RUN command is changed, the layer containing the updated packages will be reused. [Best practices for writing Dockerfiles](#) from Docker say:

When processing a RUN apt-get -y update command, the files updated in the container aren't examined to determine if a cache hit exists. In that case, just the command string itself is used to find a match.

This is acceptable if there are no other packages that need to be updated, but otherwise you can end up with old and insecure packages.

Fortunately the docker build command includes the --no-cache flag to discard the cache when building the image. This can increase the build time, so it is a tradeoff that needs to be considered when building the images.

A reasonable policy would be something like:

- Use cache on regular builds (if that happens a few times a day in CI environment, for example) to speed up the build process.
- On a regular basis (daily, weekly, on the event of a major CVE, etc.), invalidate the cache (using the –pull and –no-cache flags).

In conclusion, image build caching can be powerful. It provides significant benefits in terms of speed and efficiency when building container images, particularly in large-scale production environments. However, image build caching also introduces potential risks.

One of the main risks is that outdated or vulnerable packages or dependencies may be cached, which can lead to security vulnerabilities in the resulting images. This is especially true in environments where images are built and deployed quickly, with little time for thorough testing and security reviews.

Similarly, image build caching can lead to inconsistency in the resulting images. If an image is built with cached layers that are not up-to-date, it may not behave as expected or may not be compatible with other components in the system. This can lead to unexpected errors, downtime, or other issues.

Using admission controllers

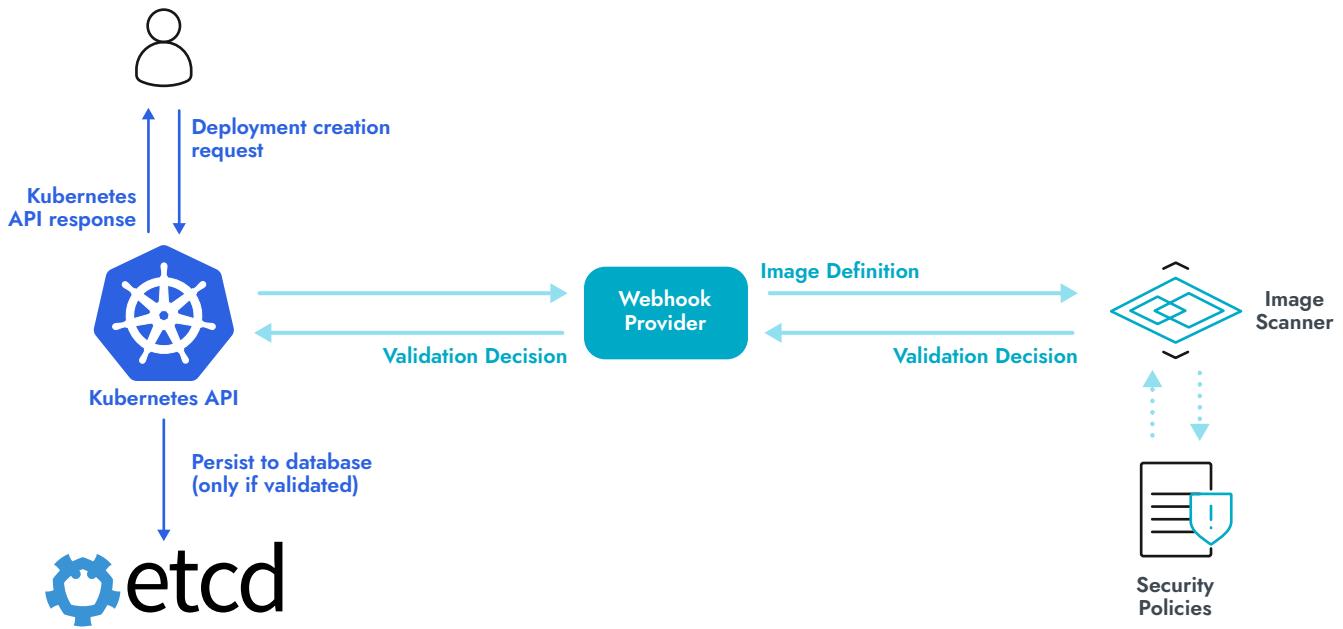
Kubernetes is designed to provide a secure and scalable environment for deploying and managing containerized applications. One of its key features is the ability to prevent unauthorized software from running within the cluster. Kubernetes achieves this by using a combination of access control policies, network isolation, and container runtime security features. These measures help ensure that only approved software is deployed, reducing the risk of malicious actors gaining access to your environment.

In addition, Kubernetes supports automation and cloud-native patterns, making it easier to manage and deploy applications at scale. This is in contrast to traditional software orchestration methods, which typically rely on manual configuration and maintenance. Kubernetes supports automation through its declarative configuration model and API-driven control plane, which enable you to define the desired state of your environment and have Kubernetes automatically manage the deployment and scaling of your applications to meet that state.

The automation capabilities of Kubernetes make it particularly well-suited for cloud-native environments, where applications are designed to be distributed, fault-tolerant, and scalable. By taking advantage of Kubernetes' features for preventing unauthorized software from running and automating key processes, you can ensure that your applications run smoothly and securely, while also reducing the amount of manual labor required to manage your environment.

Even if you identify vulnerable images in your CI/CD pipelines, often nothing is stopping them from being deployed in production. Admission control is a Kubernetes mechanism for preventing workloads from deploying to clusters unless they comply with policy, and the policy can be related to the container image's vulnerability status. Ideally, you would like Kubernetes to check the images before scheduling them, blocking unscanned or vulnerable images from being deployed onto the cluster, especially a production cluster.

[Kubernetes admission controllers](#) are a powerful Kubernetes-native feature that helps you define and customize what is allowed to run on your cluster. An admission controller intercepts and processes requests to the Kubernetes API after the request is authenticated and authorized, but prior to the persistence of the object.



Scanning tools usually offer a validating webhook that can trigger an image scanning on demand and then return a validation decision.

An admission controller can call this webhook before scheduling an image. The security validation decision returned by the webhook will be propagated back to the API server, which will reply to the original requester and only persist the object in the ETCD database if the image passed the checks.

However, the decision is made by the image scanner without any context on what is happening in the cluster. You could improve this solution by using OPA.

Open Policy Agent (OPA) is an open source and [general purpose policy engine](#) that uses a high-level declarative language called Rego. One of the key ideas behind OPA is decoupling decision making from policy enforcement.

With OPA, you can [make the admission decision in the Kubernetes cluster](#) instead of the image scanner. This way, you can use cluster information in the decision making, like namespaces, pod metadata, etc. An example would be having one policy for the “dev” namespace with more permissive rules, and then another very restrictive policy for “production.”

Open Policy Agent (OPA)

[OPA](#) is an open source policy engine that allows teams to define and enforce policies across their cloud-native environment. With OPA, teams can create policies to detect and respond to malicious or vulnerable images within their containerized applications. These policies can be customized to fit the specific needs of a team and can include rules that identify specific vulnerabilities, block malicious images, or trigger alerts.

Here is an example of a policy using OPA that blocks images with known vulnerabilities:

```
package main

import data.vulnerabilities

deny[msg] {
    input.image.vulnerabilities[_].id == vulnerabilities[_].id
    input.image.vulnerabilities[_].severity >= vulnerabilities[_].
severity
    msg := sprintf("image contains high severity vulnerability %v (%v)",
[vulnerabilities[_].id, vulnerabilities[_].severity])
}
```

In this example, the policy defines a deny rule that triggers when an image contains a vulnerability with a severity level equal to or greater than a predefined threshold. The policy uses a data file called vulnerabilities that contains a list of known vulnerabilities and their associated severity levels.

Teams can integrate this policy into their continuous integration and deployment (CI/CD) pipeline to automatically block images with known vulnerabilities from being deployed. This helps ensure that only secure images are used in production environments.

OPA Gatekeeper

Tools like OPA and Falco can help teams respond to vulnerabilities and malicious images by defining and enforcing policies across your cloud-native environment. In the case of Kubernetes, [OPA Gatekeeper](#) is an additional policy engine that could be used to enforce policies on K8s-specific resources.

With Gatekeeper, teams can define policies that enforce specific configurations, security requirements, or other constraints on Kubernetes resources. Here's an example policy for detecting vulnerable images using Gatekeeper:

```
package k8srequiredlabels
import data.k8srequiredlabels
violation[{"msg": msg}] {
    image := input.review.object.spec.template.spec.containers[_].image
    vulnerabilities := data.vulnerabilities[image]
    vulnerabilities != null
    vulnerabilities[_].severity >= 7
    msg := sprintf("Vulnerable image %s found in container %s", [image, ])
```

The above policy checks the severity level of vulnerabilities for a given container image, and if it is equal to or greater than 7 (on a scale of 0 to 10, for example), it generates a violation message indicating that a vulnerable image has been found in a container. This policy assumes that a data source called "vulnerabilities" is provided that maps image names to vulnerability information.

Constraints are defined using OPA Rego language and can be customized to meet the specific needs of an organization or application. They can be used to enforce a wide range of policies, such as requiring all pods to run with specific security settings or, in this case, disallowing the use of certain container images.

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: vulnerable-image
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Deployment"]
  parameters:
    labels:
      requiredLabel: "true"
  enforcementAction: deny
  audit:
    namespaces:
      - default
  rego:
    content: |
      package k8srequiredlabels
      import data.vulnerabilities
      violation[{"msg": msg}] {
        image := input.review.object.spec.template.spec.containers[_].
image
        vulnerabilities := data.vulnerabilities[image]
        vulnerabilities != null
        vulnerabilities[_].severity >= 7
        msg := sprintf("Vulnerable image %s found in container %s",
[image, input.review.object.metadata.name])
      }
```

This constraint uses the policy defined earlier to deny deployments that use vulnerable images. It matches on Deployments, audits the “default” namespace, and enforces the policy by denying non-compliant resources. With this constraint in place, Gatekeeper will monitor new deployments and prevent vulnerable images from being used at runtime.

Dedicated scanning tools

While image scanning is a common feature, it is typically considered part of SCA (Software Composition Analysis), which is in turn part of AST (Application Security Testing). To address the unique challenges of vulnerability management in cloud-native architectures, it is important to use tools and practices that are specifically designed for these environments. This may include SCA tools that can scan container images for vulnerabilities, as well as integration with orchestration platforms like Kubernetes to automate vulnerability scanning and patching processes. By taking advantage of these dedicated tools, you can more effectively manage vulnerabilities in your cloud-native environment and reduce the risk of security breaches.

As the demand for cloud-compatible vulnerability management tools grows, traditional vendors are expanding their offerings to include solutions that can scan IaaS, containers, and other cloud-native environments. However, it's important to note that these tools may have limitations when it comes to cloud-specific features, such as scan speed and visibility into serverless functions. In other words, while these tools may be adapted for the cloud, they may not be optimized for it, and may not be as effective as purpose-built cloud-native vulnerability management solutions.

Some of the open source tools available include:

- [Trivy](#): A comprehensive and versatile security scanner.
- [Clair](#): An app for parsing image contents and reporting vulnerabilities affecting the contents.

Additionally, commercial options can be licensed to perform cloud-native vulnerability management at scale. You can learn more about Sysdig Secure [here](#).

Summary/Conclusion

As businesses continue to adopt cloud-native architectures, the need for effective vulnerability management becomes increasingly important. The dynamic and distributed nature of these architectures require a new approach to vulnerability management that is tailored to the unique challenges of the cloud.

As a result, an effective vulnerability management strategy for cloud-native architectures requires a shift in mindset from traditional vulnerability management practices. It involves continuous monitoring and assessment of the entire cloud environment, including containers, microservices, and serverless functions. It also requires automation and collaboration between development, security, and operations teams.

Sysdig enables prioritization with runtime context throughout the cloud innovation lifecycle, addressing pain points for DevOps and Security teams. This includes establishing Falco detection as a cornerstone for real-time runtime insights.

By adopting a proactive approach to vulnerability management, using tools like Falco and OPA, organizations can reduce the risk of cyberattacks, data breaches, and other security incidents. We highlight some real-world breaches where applying vulnerability scanning capability could have prevented the breach, therefore ensuring compliance with industry regulations and standards is met, as well as building trust with their customers and stakeholders.

In summary, evolving vulnerability management for cloud-native architectures is a critical step towards ensuring the security and reliability of cloud-based systems. As the cloud continues to play an increasingly important role in modern business, organizations must prioritize their vulnerability management strategies to stay ahead of potential threats and protect their sensitive data and assets.

Stay informed and ahead of the curve with the Sysdig 2023 cloud usage and security report, which provides valuable insights into the latest threats and trends in cloud-native environments. You can access the report at <https://sysdig.com/2022-cloud-native-security-and-usage-report/> to gain a deeper understanding of the evolving landscape and take proactive steps to secure your environment. Don't wait until it's too late — stay ahead of the game with Sysdig.

Powered by runtime insights, Sysdig stops threats instantly and reduces vulnerabilities by up to 95%. We created Falco, the open standard for cloud threat detection, and apply runtime insights to help you focus on the vulnerabilities and threats that matter most. Prevent, detect, and respond at cloud speed with Sysdig.

REQUEST A DEMO

