

RSA ConferenceTM 2024

San Francisco | May 6 – 9 | Moscone Center

SESSION ID: SBX-R06

Kubernetes Security: Attacking And Defending Modern Infrastructure

Lenin Alevski

Security Engineer
Google
@alevsk

THE ART OF
POSSIBLE



#RSAC

Max vonBlankenburg

Security Researcher
Semgrep
@hackfidgetcube

Disclaimer

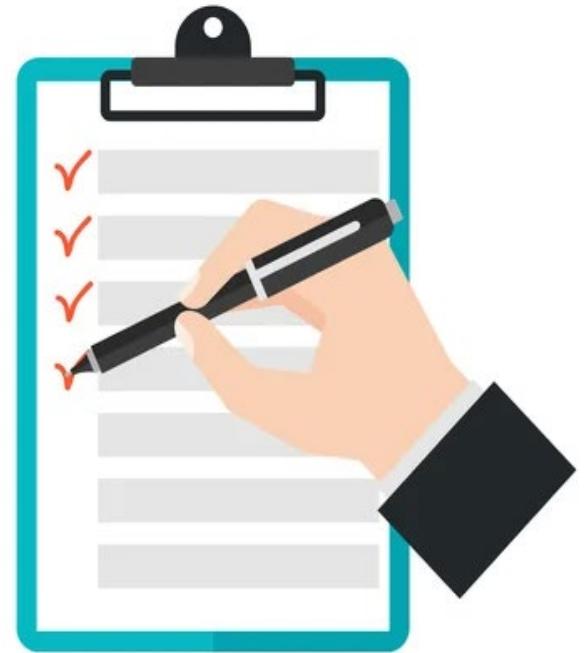
Presentations are intended for educational purposes only and do not replace independent professional judgment. Statements of fact and opinions expressed are those of the presenters individually and, unless expressly stated to the contrary, are not the opinion or position of RSA Conference™ or any other co-sponsors. RSA Conference does not endorse or approve, and assumes no responsibility for, the content, accuracy or completeness of the information presented.

Attendees should note that sessions may be audio- or video-recorded and may be published in various media, including print, audio and video formats without further notice. The presentation template and any media capture are subject to copyright protection.

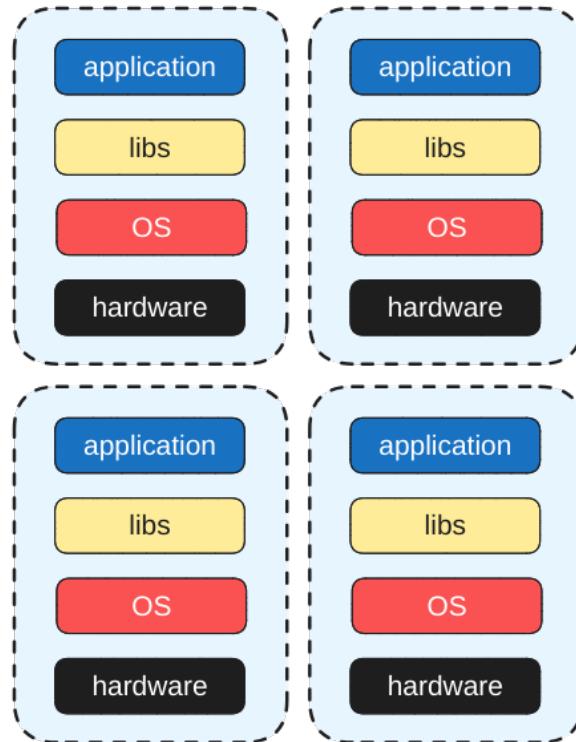
© 2024 RSA Conference LLC or its affiliates. The RSA Conference logo and other trademarks are proprietary. All rights reserved.

Agenda

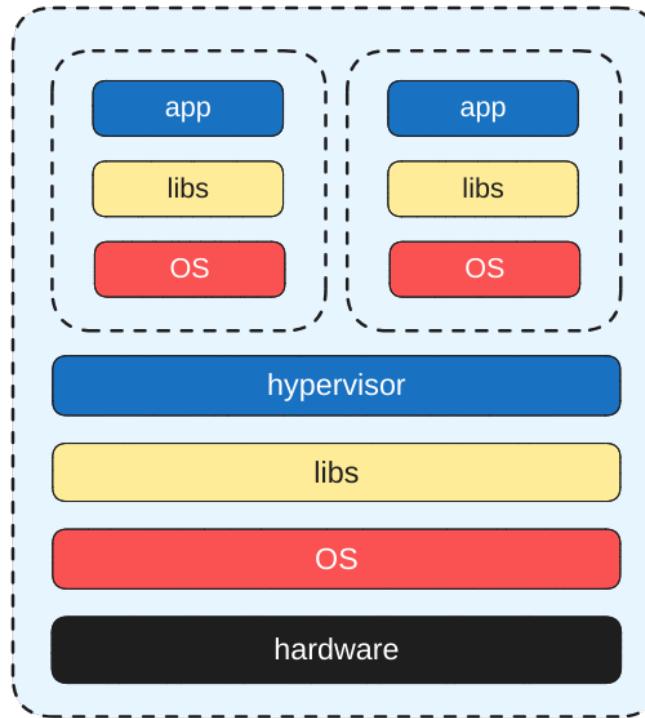
- Introduction to Containers
- Introduction to Kubernetes
- Kubernetes Threat Model
- Most common attack techniques in K8S
- Build-In Defenses
- Kubernetes Evolution



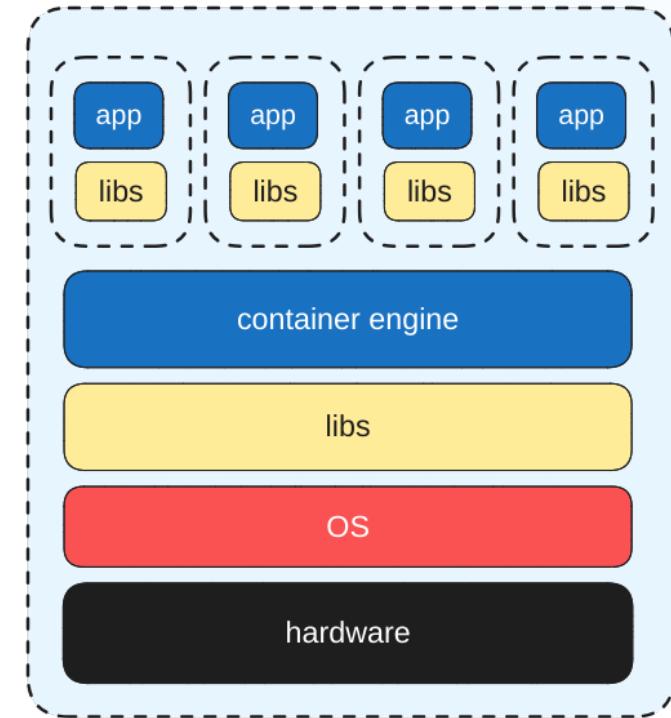
Application Compartmentalization



**Traditional
Architecture**

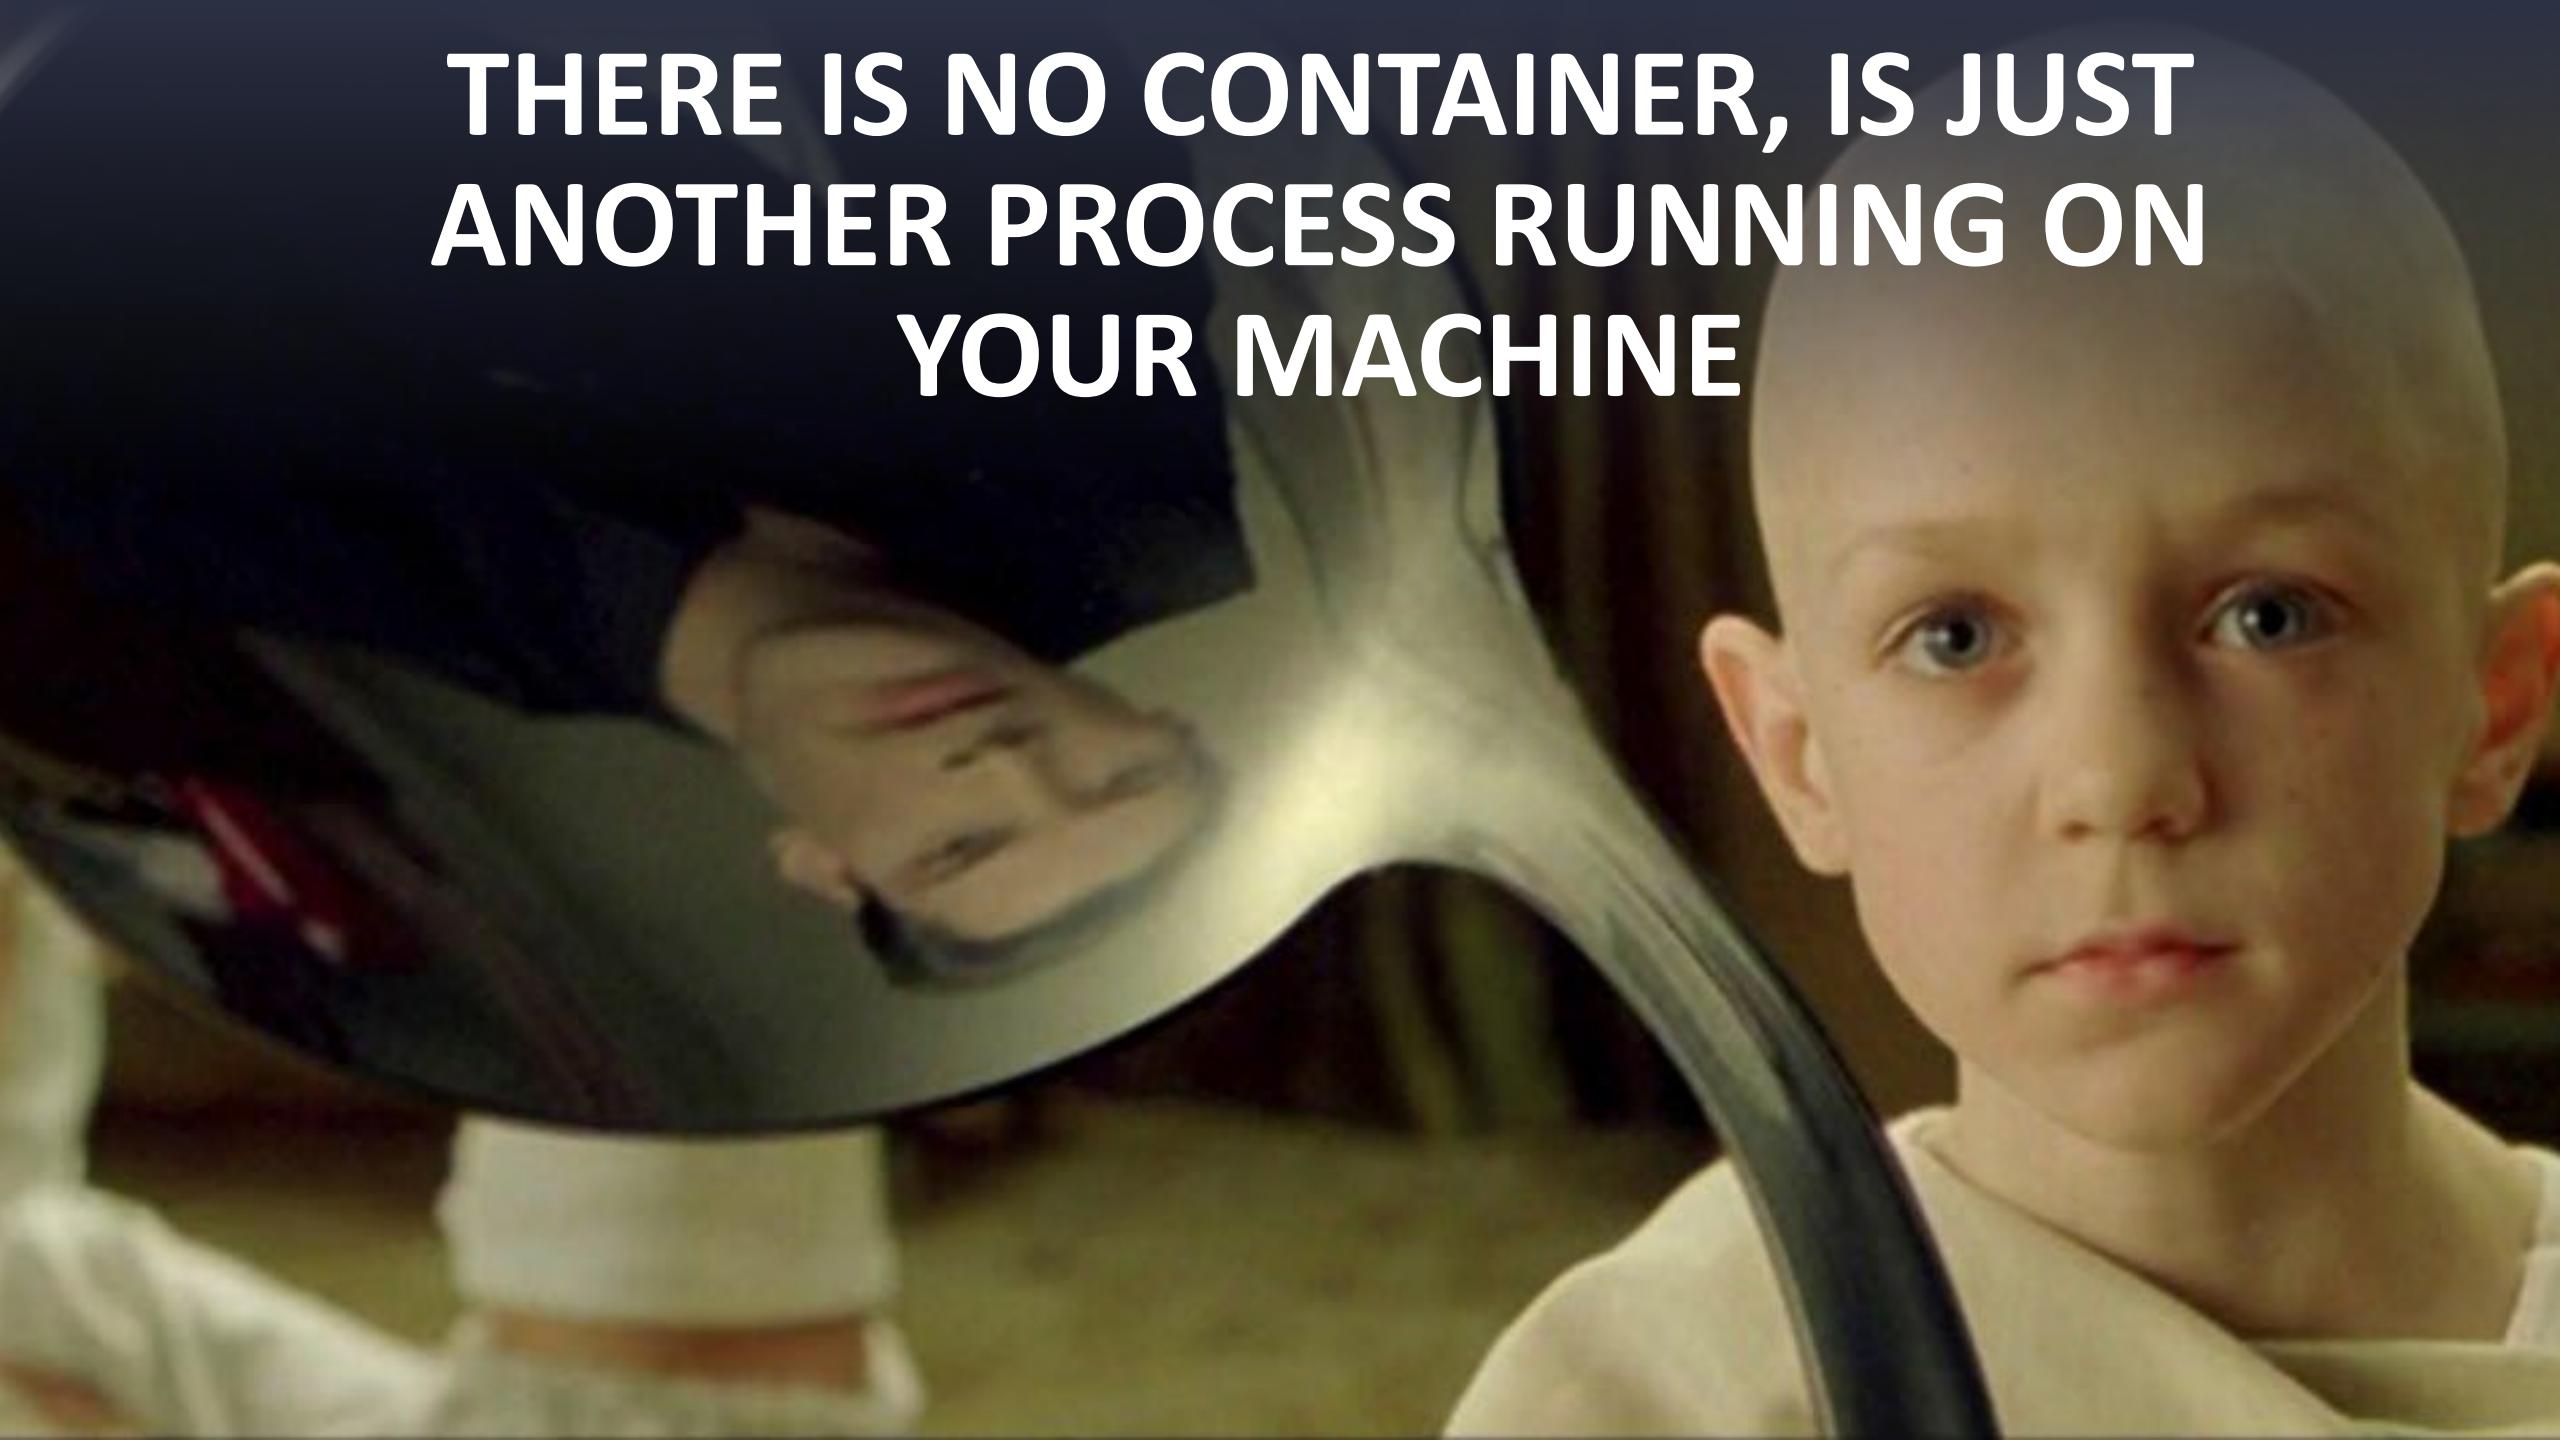


**VM
Architecture**



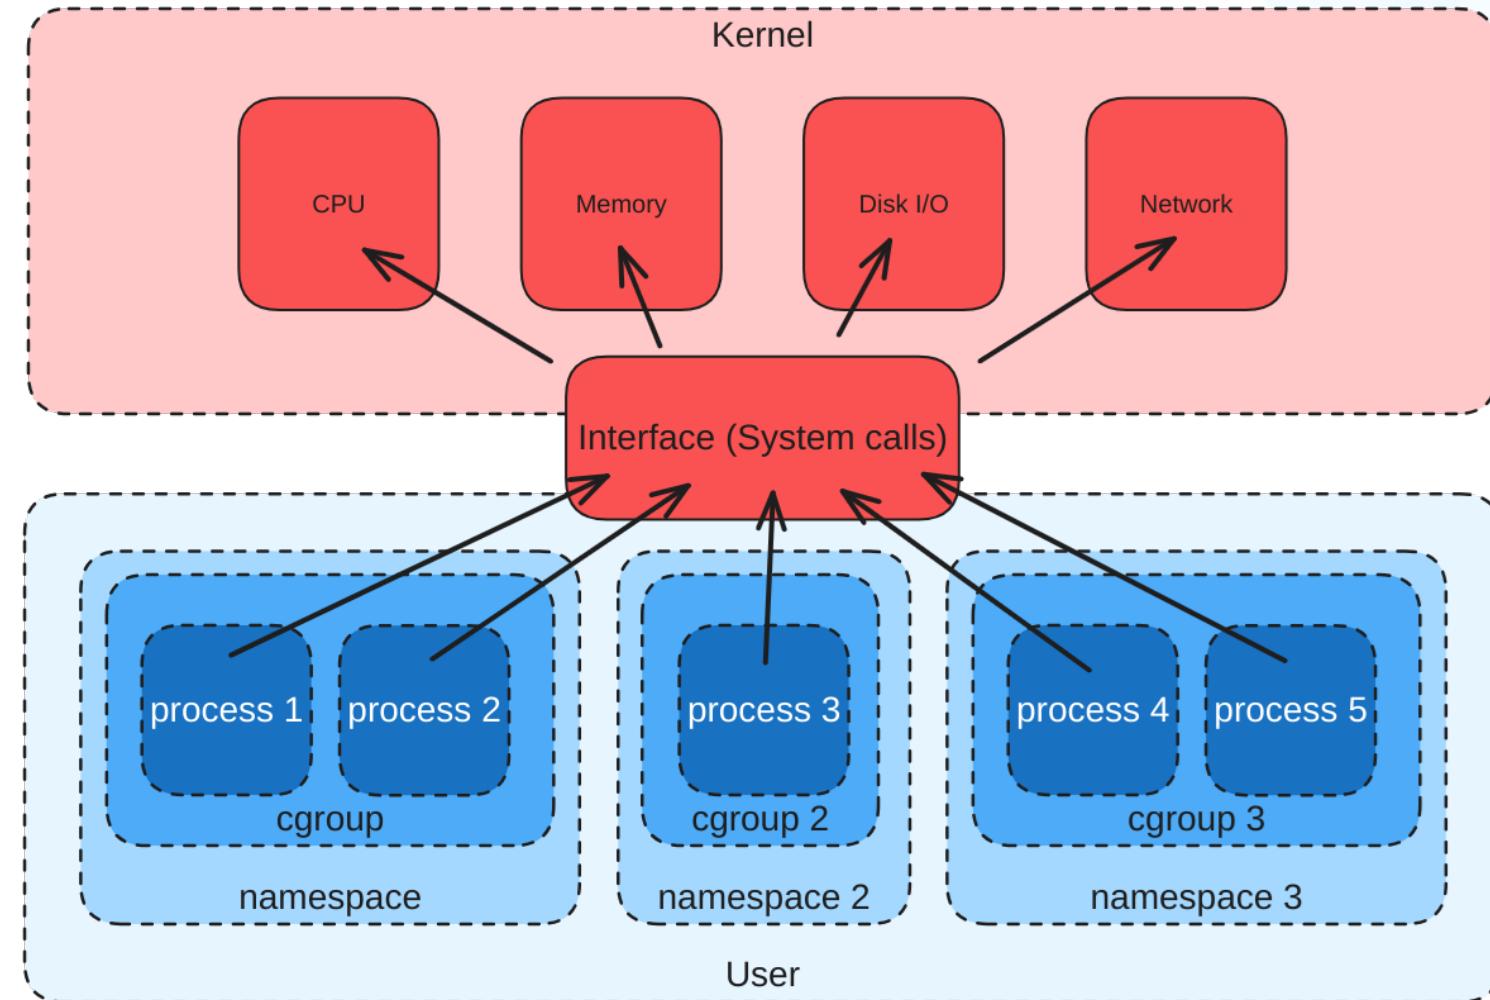
**Container
Architecture**

THERE IS NO CONTAINER, IS JUST
ANOTHER PROCESS RUNNING ON
YOUR MACHINE

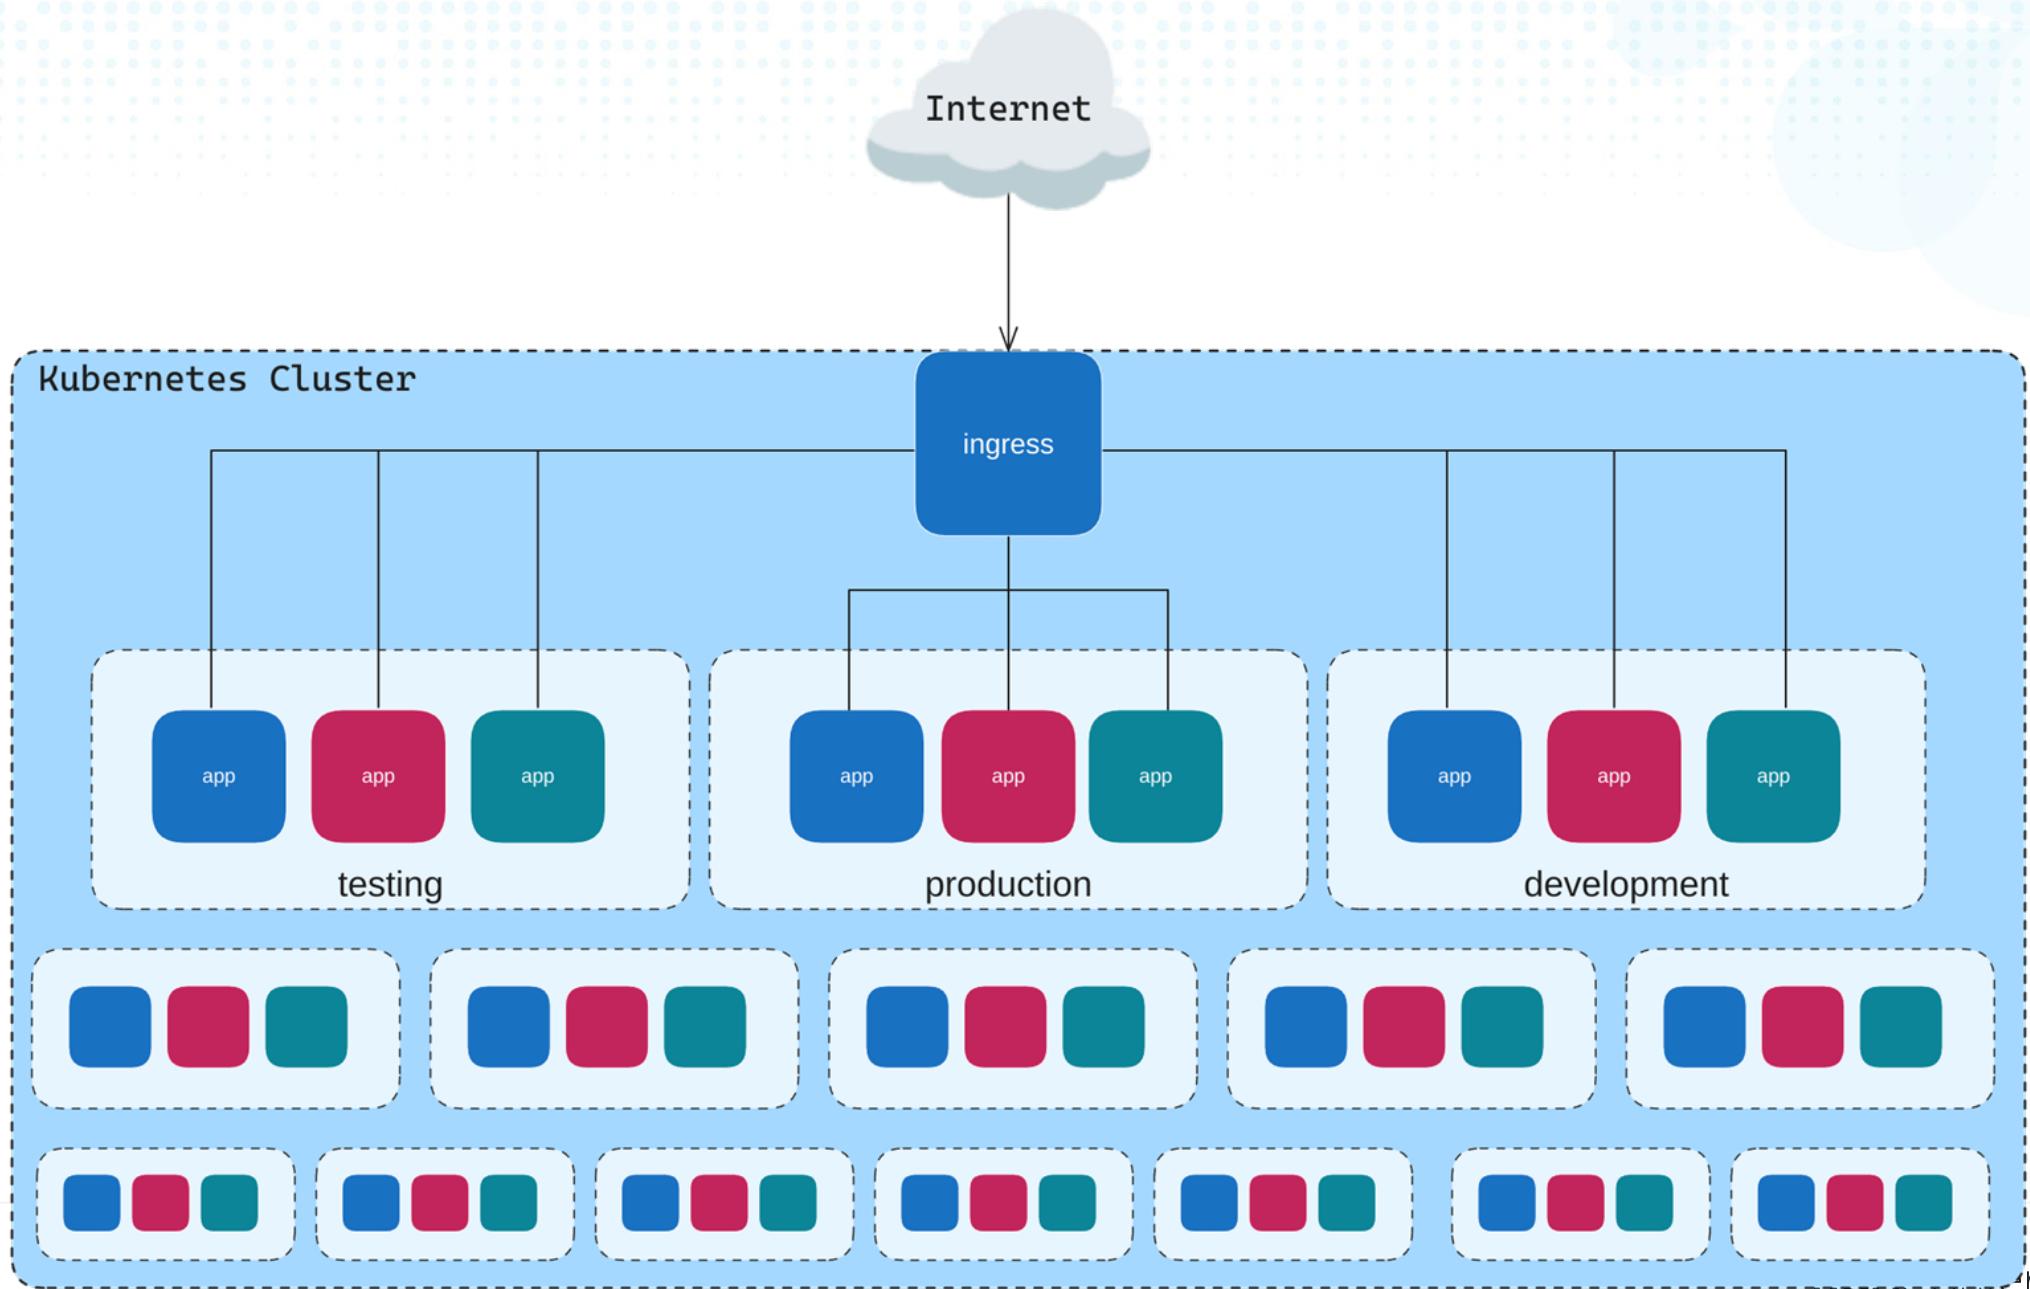


So, What Containers Really Are?

- Namespaces
- Cgroups
- Capabilities



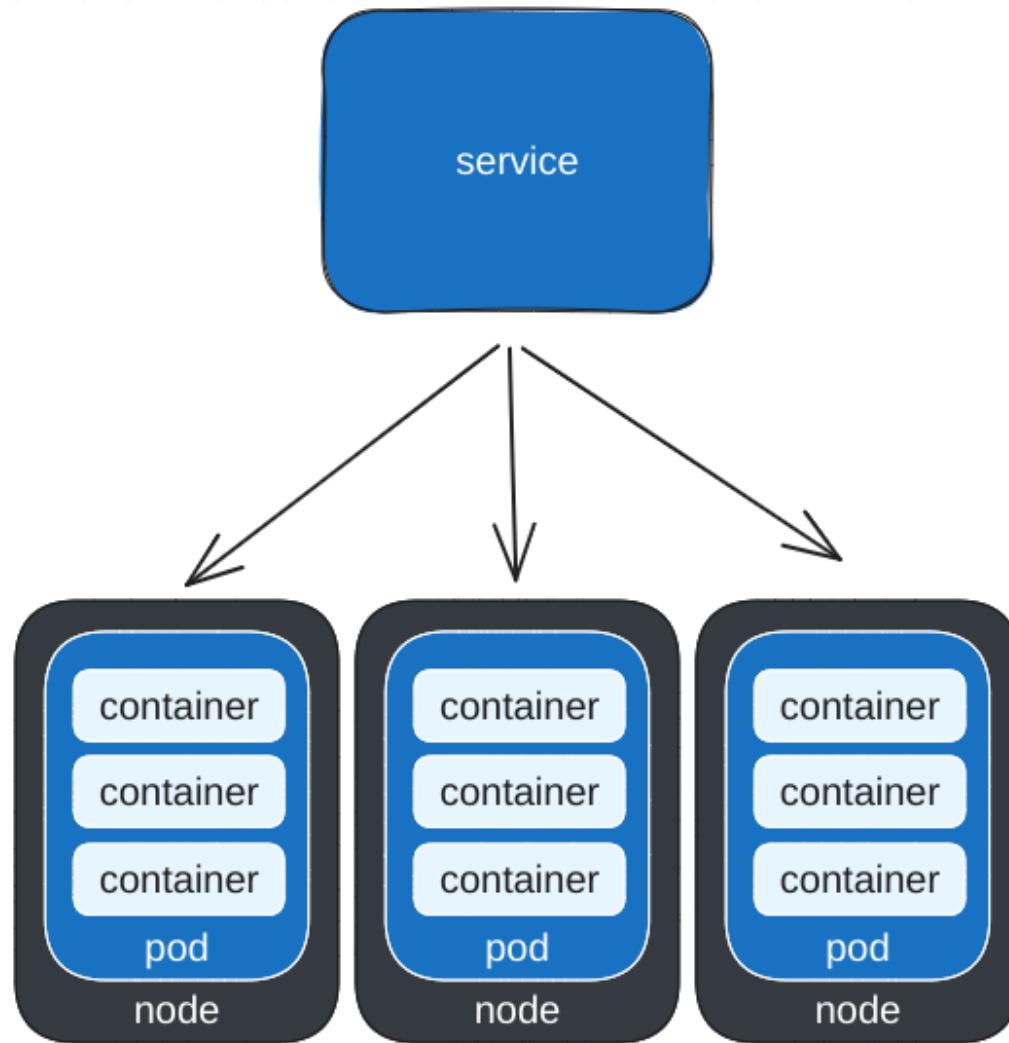
What Is Kubernetes Anyways?



Kubernetes Primitives

Ingress LimitRange
ConfigMap ComponentStatus
PersistentVolume NetworkPolicy
DaemonSet ClusterRole
Service Namespace
clusterRoleBinding
Job Volume
HorizontalPodAutoscaler Node Pod Role
CronJob Secret
(CRD) ServiceAccount
ServiceAccount EventPriorityClass

Kubernetes Application



```
apiVersion: v1
kind: Service
metadata:
  name: my-app
spec:
  selector:
  ...
---
apiVersion: apps/v1
kind: Deployment
metadata:
  ...
spec:
  ...
  replicas: 3
  template:
  ...
  spec:
    containers:
    ...
    ...
    ...
```

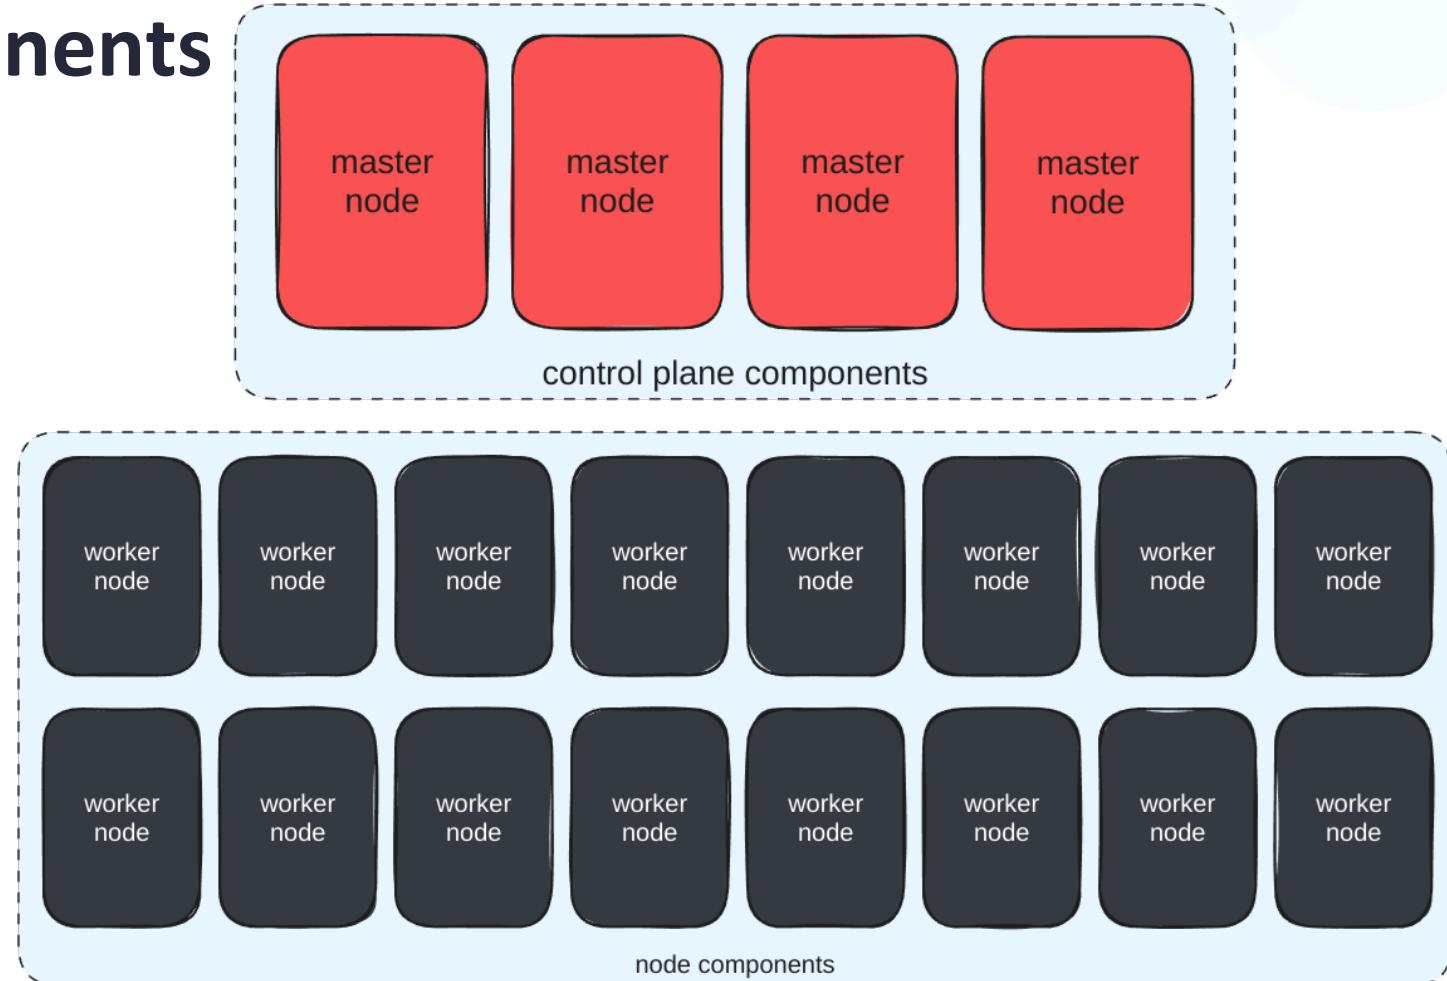
Kubernetes Components

Control Plane Components

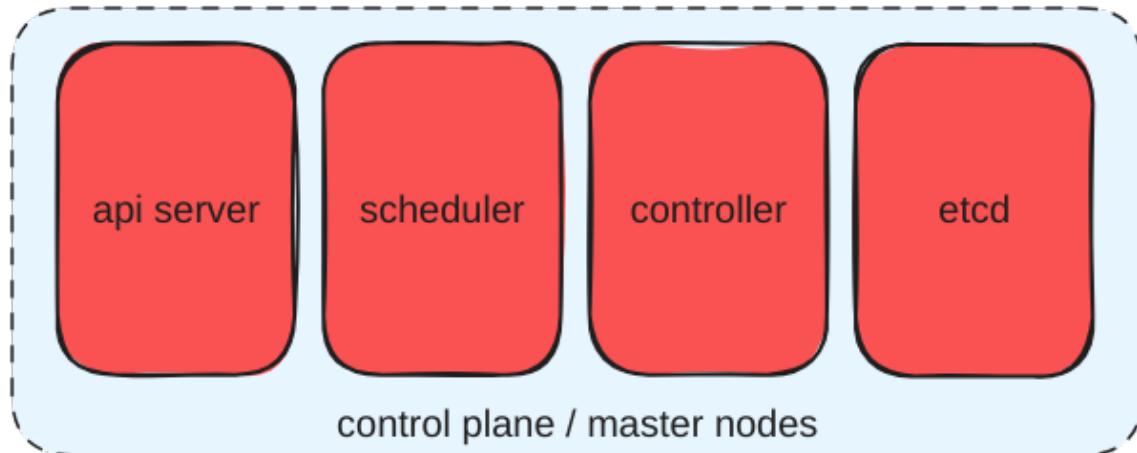
- kube-apiserver
- etcd
- kube-scheduler
- kube-controller

Node Components

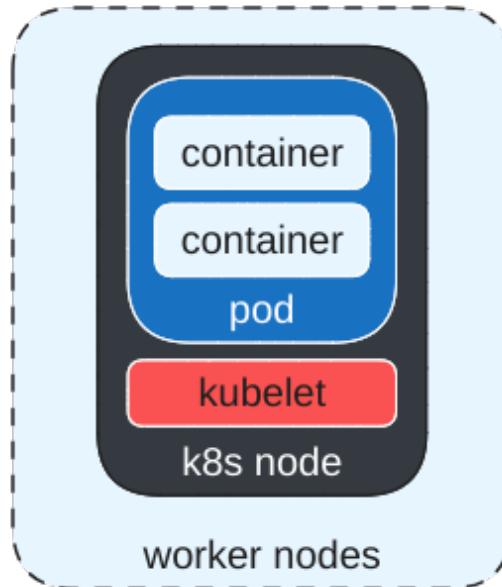
- kube-proxy
- kubelet
- container runtime



Kubernetes Threat Model: If An Attacker Controls...



Control plane nodes: Attacker controls your cluster. They can modify, access and destroy everything



Pod / Container: Attacker controls application and may be able to escape and attack the node

Kubelet: Attacker controls running pods

Worker nodes: Attacker controls running pods. They can attack master nodes

OWASP Kubernetes Risks Top (2022)

- K01: Insecure Workload Configurations
- K02: Supply Chain Vulnerabilities
- K03: Overly Permissive RBAC Configurations
- K04: Lack of Centralized Policy Enforcement
- K05: Inadequate Logging and Monitoring
- K06: Broken Authentication Mechanisms
- K07: Missing Network Segmentation Controls
- K08: Secrets Management Failures
- K09: Misconfigured Cluster Components
- K10: Outdated and Vulnerable Kubernetes Components

Kubernetes Most Common Attack Techniques

Threat Matrix For Kubernetes (2020)

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account	Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking	Denial of service
Application vulnerability	Application exploit (RCE)		Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files	
Exposed Dashboard	SSH server running inside container					Instance Metadata API	Writable volume mounts on the host	
							Access Kubernetes dashboard	
							Access tiller endpoint	

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Images from a private registry	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account		Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files		
Exposed Dashboard	SSH server running inside container				Access managed identity credential	Instance Metadata API	Writable volume mounts on the host		
Exposed sensitive interfaces	Sidecar injection				Malicious admission controller			Access Kubernetes dashboard	
								Access tiller endpoint	
								CoreDNS poisoning	
								ARP poisoning and IP spoofing	

Threat Matrix For Kubernetes (2021)

= New technique

= Deprecated technique

Initial Access

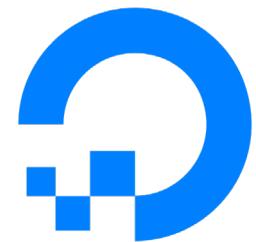
- Using cloud credentials
- Compromised images / registry
- Kubeconfig file
- Application Vulnerability
- Exposed sensitive interfaces



Using Cloud Credentials



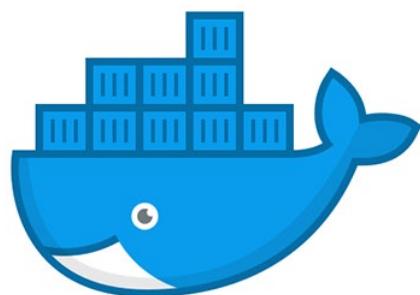
Google Cloud



DigitalOcean

Compromised Registry And Images

- Supply Chain Attacks
- Vulnerable dependencies on images



Kubeconfig File - Hunting For .kube/config Files

```

apiVersion: v1
clusters:
- cluster:
  certificate-authority-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM2akNDQwRLZ0F3SJUJB01C0VRBTkJna3Foa2lHOxrw0kFrC0ZBREftTVNRd0lnwURWUVFEREJ0dmNHVnUKYzJocFpuUXrjMmxuYm1WeVf
UREoN2Y0dwdMyahBab1FY0zJsbm1tvnLREUw70RnE1U1UTB0VGN3CndrRwLN0TBH01Nx1Nj1neUUV0V0FVQUE0SU1Ed0f3Z2dFS0fvsLUJBURML2deBfNmZnExK3NZUUnnTz2ZEYKE0EVs
NPTW91WFdKwnk5NVPQ2uWpEgGb3kyTUNnCs5ZkXVNE19nUE52c0RYSTFs0t1RTBnFzZaaHjHckFTNHjYrJxdLjwRkhJ03vtVy9nYUFYnJla2Ns5d0N1RxBqM3hnL0pwamVuV2NvVHBU0pS
i1Nz1nq0V0m05d01d1hnZ0J1M1vU5G5tC9K3psm53NkdsxaRCUUR1Tx0c2hNxEttUmprU6T6RGL1XTUcxM1zRCKKfnTuJ3B0ldqgXpbAe1BNedMVkRhfdF0193UUV8d010cERBUjeJnTU
NU4Udm16VUyjbx1B0WV3UiMagpP0kVksTNGZLfuUD21mNFej11em0vNlqvNzdMajEvTHkRmpGMjV2Edjvms5Yx0cL55V1ybnnM4QlxNT0FBcnZkVXSRwpvXwJU0Jm3eGp1Vmxi.cVzeDF
Gjckx4cn5xd0ZqRFZ3MURw5214wDznwU9V0wRxwyp6ZkpByLhbxLkvT1kRkMwRTYza12hRy91RwZudFBhJzdB3hJzLzxaHF321RhzMk50uubkrMmY4ZCtsZLm4CjAzelF0Nmh3jh
server: https://10.111.112.101:8443
name: 10-111-112-101:8443
contexts:
- context:
  cluster: 10-111-112-101:8443
  user: scruffy/10-111-112-101:8443
name: /10-111-112-101:8443/scruffy
- context:
  cluster: 10-111-112-101:8443
  namespace: default
  user: system:admin/10-111-112-101:8443
name: default/10-111-112-101:8443/system:admin
- context:
  cluster: 10-111-112-101:8443
  namespace: marmalade
  user: scruffy/10-111-112-101:8443
name: marmalade/10-111-112-101:8443/scruffy
- context:
  cluster: 10-111-112-101:8443
  namespace: pizzazz
  user: scruffy/10-111-112-101:8443
name: pizzazz/10-111-112-101:8443/scruffy
current-context: default/10-111-112-101:8443/system:admin
kind: Config
preferences: {}
users:
  name: scruffy/10-111-112-101:8443
  user:
    token: RmIp2WbgoiYM8BNNxtodGUk3m1xjpK3jntHvq8yfrs
  - name: system:admin/10-111-112-101:8443
    user:
      client-certificate-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURENDQwZX0F3SJUJB01C0VRBTkJna3Foa2lHOxrw0kFrC0ZBREftTVNRd0lnwURWUVFEREJ0dmNHVnUKYzJocFpuUXrjMmxuYm1WeVf
UtFeF262vH0MpxMDZMmNgYzNs6GnpMhArzFvN5neEZUv0RCZ05WCKj3BTvRESE41YzNsK6jGhaRzFwYmpd0Q9FTSx0dEU1KS29aSvh2Y05BUvCQ1FBRGdnRVBBREN0Vf02dnRUkQ01z
c4ZGeYVCvQx0u3NkbXm0Lu1v0t0vH05yNGLjMnVrteZMRTBVFDf2MGFcqdhaEk0bmdoaVNFnzhPcjhBm19xWlLDT19BCNvTFhLmnV0NmhdEfwaU8200FHu2JjQ0F3RUFByU0xTURN0dRnMu
PvkD0wprpRHN2d0J3de9abFnryFyJw05NGLjMnVrteZMRTBVFDf2MGFcqdhaEk0bmdoaVNFnzhPcjhBm19xWlLDT19BCNvTFhLmnV0NmhdEfwaU8200FHu2JjQ0F3RUFByU0xTURN0dRnMu
b1pjahZ1TkFRRUxCUFEz2dF0gBSE1oMD10WjFyRo1beFwy2g30UvFzUf0qenNxV1kQ1UrUwVM6Xu30dssnVfsj1PRW9mt0c5zVG0GpbdFxChnVsjhjdzB2zg2MldJ0tVCREJGzJRCRC
Vnm0HNTuXFMwWZ3dnK0NMHh1b2Y5T5kneXhyjXEV0q0LaGjW5TyUmszTUDcN2LTWq0k2tU03n1lUmdWkhkWvA4a0lPwmhD05TzSwxZng1qGF0UnFwRjhaK1dZa2p0CjNzU3EP03pIV3B
g0Mz2vM5M5C101v50tRUE1ENFU1RjkLdQVRFLS0tLS0K
  - client-key-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURENDQwZX0F3SJUJB01C0VRBTkJna3Foa2lHOxrw0kFrC0ZBREftTVNRd0lnwURWUVFEREJ0dmNHVnUKYzJocFpuUXrjMmxuYm1WeVf
0xEMj1tX13aHe0M1pnadhp0XZKdx1Yhd2b1FYy18x0E2skNzZ5HkmndBmpwX0HE4NH1kdnReB6jzZuN0lGvMtjh1sld0QVYK0tNwk1D0k1STbVh4z2Rhd1xysEpsvTdrwE1JRUs32z1V4Vfk3
dEYj1ab0dJmHbs2d022JhnC9xr1swWfNjN330vpkhdJREFRQ0B01C0VFDtxV-NXyyRn1jb1jG0Nq0pVjdyW09vNj0V0GdvTDRdk2ZxNhVhVn1TeFvW10hUoU1Eh31TTRtgdRwh1MnMzeXhDw
XynrvaU0db01K1n14j2X0V1Xewp0cmpEMEgsXk80d0vh005sPf1RnL2zRSt1FYTU1c19Yhndj0kvzThfUjpbnj0Ue14awo4Nmzw2Ez2FdWf21xYkfBzBzCgr0u3F42EBw0Wtslh0Uj
Uk14ZtL6NmNz1N1Tx2NfIem0wFHKYloak5vWhbb6dCQ19j1XRIUWjshM0VZMHN1UhwsCpWpYlFN3A0vJNMFVNV0Vw0dkpNmNzK2VHMN1Lz1Uego1M0jzcnrbG92cm5tSFRCnvUyZU
zf521r2a09P0t1vM1hnb16cgPhlm9Qn0h0HQ10W9H0QkPFPSfUKL3E-P830EFZUGovKv01W1TzHxNzVnNkSuNwczV0M1hmjB165d1uXhMmWj0Mn1wazc5SLR PangySLV6V7btwdpVSY0hjV
RntNtgYUZ1TgHyrVz01V2Rkx0J0wPj0n1zWkfVr0FJn0vEBUz6WjRSUy4UxgMrUvUhBdH12VmxCvHwODNv0Wkl2xMfJLURk0r52F2zDvU0h5pQVz2VULWjNKV0dJnU5SY04y5zdy5Pub
pcn8xLnfs1p0MEJPZnFwauJ7QwPj0n1zWkfVr0FJn0vEBUz6WjRSUy4UxgMrUvUhBdH12VmxCvHwODNv0Wkl2xMfJLURk0r52F2zDvU0h5pQVz2VULWjNKV0dJnU5SY04y5zdy5Pub
aExvUNlahESXN3a5s3Cmnb04REV0Z1C1Yk1ia3Nk1h0Bb0HfJdENFU1JH0VorevlnXVfkQkIVN0tUkrLfLWmgysDVGblhafkVkd2ckdm1s5Opqd11yekvTVWlsQ1VHUhJUrtNzktRCC
1JBRLyhJnnVku0U0LURnVqVn5Pqot50tUV0RCBSU0egUfJjkVURSBLRvtL50tQo

```

60 code results

Repositories

- Code: 60
- Commits: 0
- Issues: 18
- Discussions: 0
- Packages: 0
- Marketplace: 0
- Topics: 0
- Wikis: 1
- Users: 0

60 code results

vermegil/automate-private-AKS-AAD /AKS-AAD-RBAC-NOTES.MD

```

13 ---yaml
14 apiVersion: v1
15 clusters:
16 - cluster:
  certificate-authority-data: LS0tLS0K (this is base64 encoded certificate data
for the cluster)
...
59 - cluster:
60   certificate-authority-data: LS0..0tLS0K
61   server: https://xstofaks-79f2f498.hcp.westeurope.azure.k8s.io:443

```

Markdown Showing the top three matches Last indexed on Jan 24

Languages

Language	Count
Markdown	13
HTML	2
AsciDoc	1
YAML	19
Text	9
HAProxy	2

Advanced search Cheat sheet

kspace5/svc-mesh-demo cluster1/cluster-config.yaml

```

6   server: "https://ubsonita/k8s/clusters/c-t95zn"
7   certificate-authority-data: "LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUJpRENDQv
...
16   Uhtam54Uj1NyWorrcGhlLw1ZVGt4M3V4QuInTz10wG5UeXdcQ1TBLQVRhS3k0b0N0TGPnLNevdQp55
17   W1cE9Mb2s3VKNGb2M9Cio1Ls0tRUE1ENFu1RjkLdQvRFLs0tLS0
18 - name: "kube1-ubsonita"
19   cluster:
20   server: "https://192.168.86.68:6443"

```

YAML Showing the top two matches Last indexed on Apr 22, 2021

patiavit13/patiavit13 /kubeconfig-rancher.cfg

```

6   server: "https://203.151.50.20/k8s/clusters/c-bfhk6"
7   certificate-authority-data: "LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUJpRENDQv
...
16   TQ2dWxkQVZQhC9pcwVkMDVYNUj91MFZyWFFDSwRwGn1uvV1c3NmXR0z1M1yJ1Va1hdjLzJ0VwpQ0
17   VYveHZ5T29rZUt2dplkci0tLs0tRUE1ENFu1RjkLdQvRFLs0tLS0
18 - name: "immortal-cluster-k8s-master-2"
19   cluster:

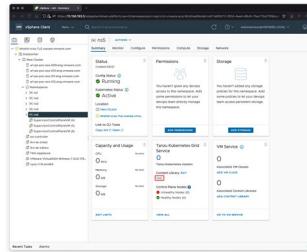
```

Application Vulnerabilities

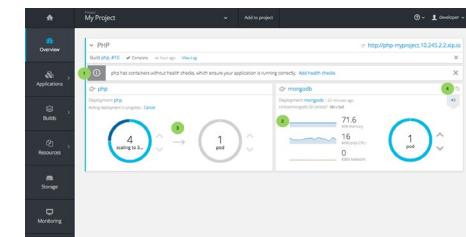
- OWASP Top 10
- SQLi
- RCE
- Command injection
- Etc

A01:2021	Broken Access Control
A02:2021	Cryptographic Failures
A03:2021	Injection
A04:2021	Insecure Design
A05:2021	Security Misconfiguration
A06:2021	Vulnerable and Outdated Components
A07:2021	Identification and Authentication Failures
A08:2021	Software and Data Integrity Failures
A09:2021	Security Logging and Monitoring Failures
A10:2021	Server-Side Request Forgery

Exposed Sensitive Interfaces

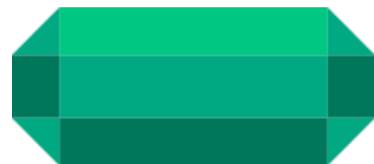


VMware Tanzu



Create New K8s Tenant

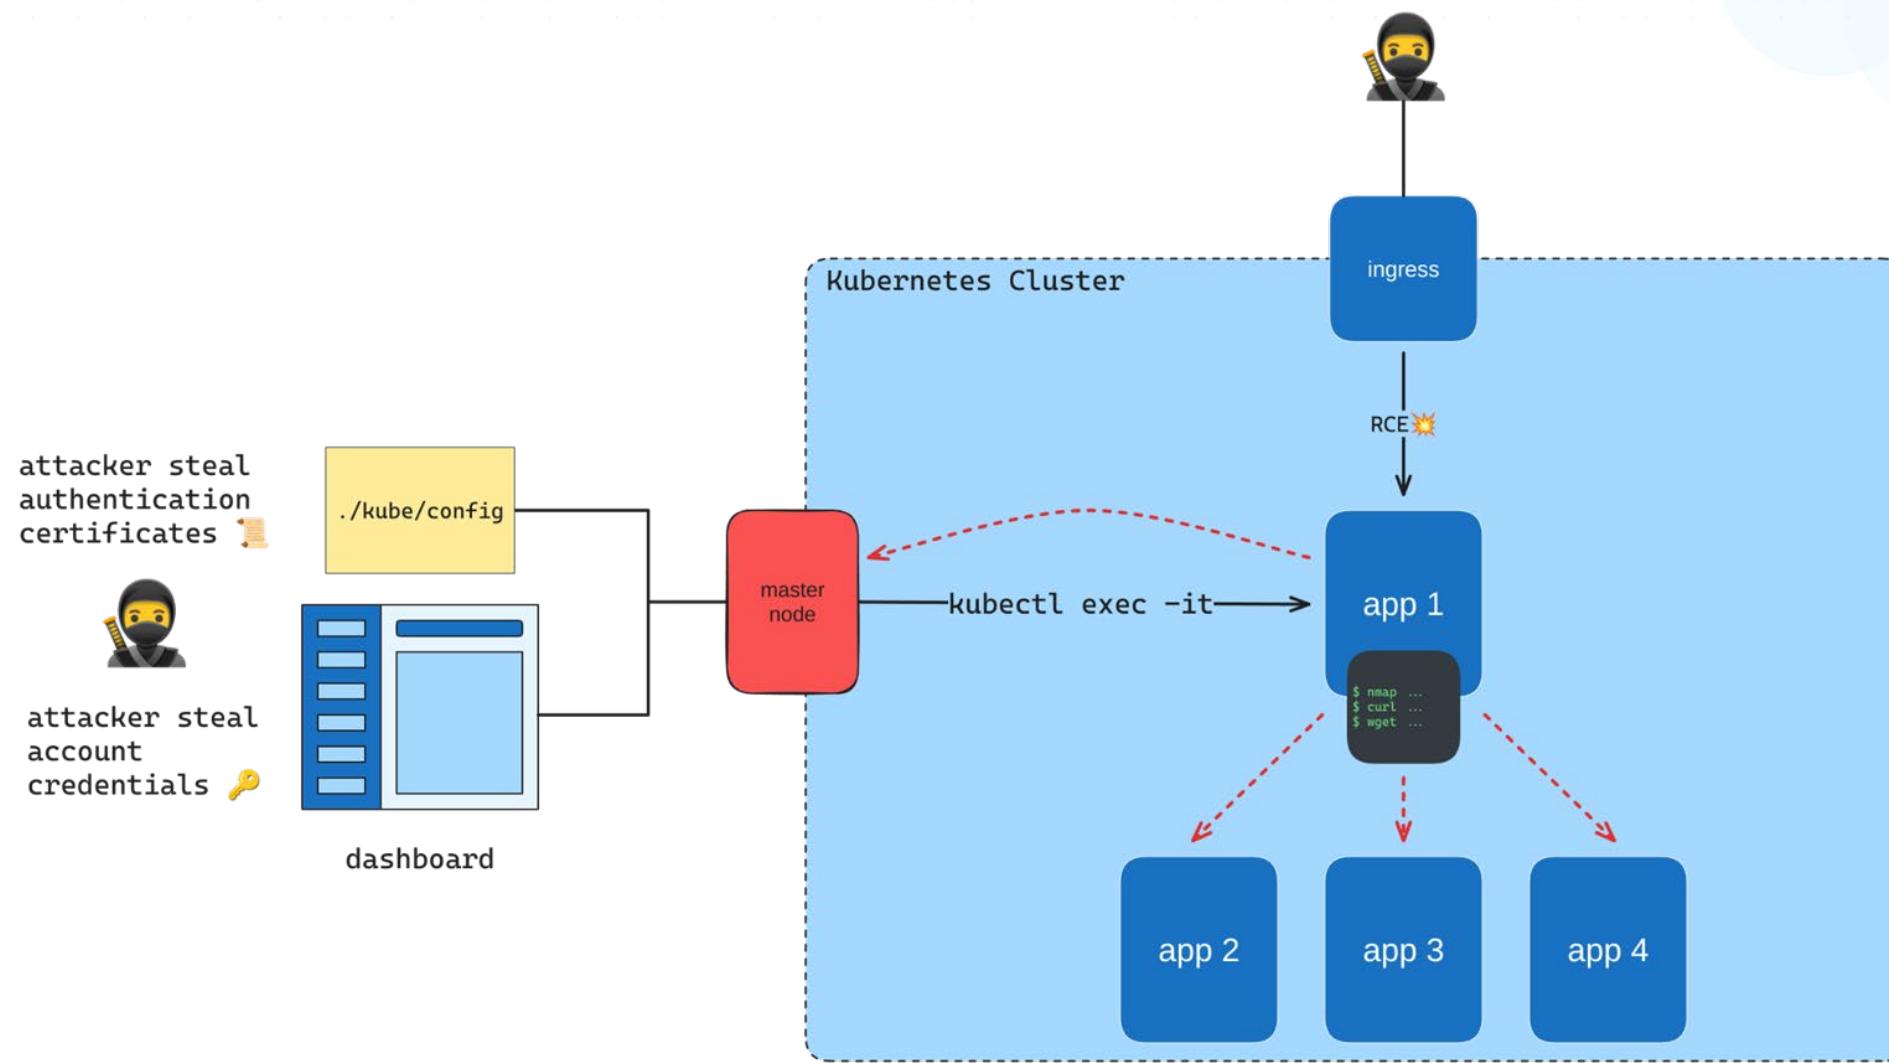
Tenant Name	<input type="text"/>
Tenant Description	<input type="text"/>
K8s Cluster	<input type="text"/> cluster1
Adopt Existing Namespace	<input type="checkbox"/> No free namespaces available to adopt in this cluster.
Specified Namespace Name	<input type="text"/>
Is Namespace Owner	<input checked="" type="checkbox"/>
Map Services To Gateway	<input checked="" type="checkbox"/> (optional)
Enable Intra Service Mesh	<input checked="" type="checkbox"/> (optional)
AWS Lambda Project	<input type="checkbox"/>



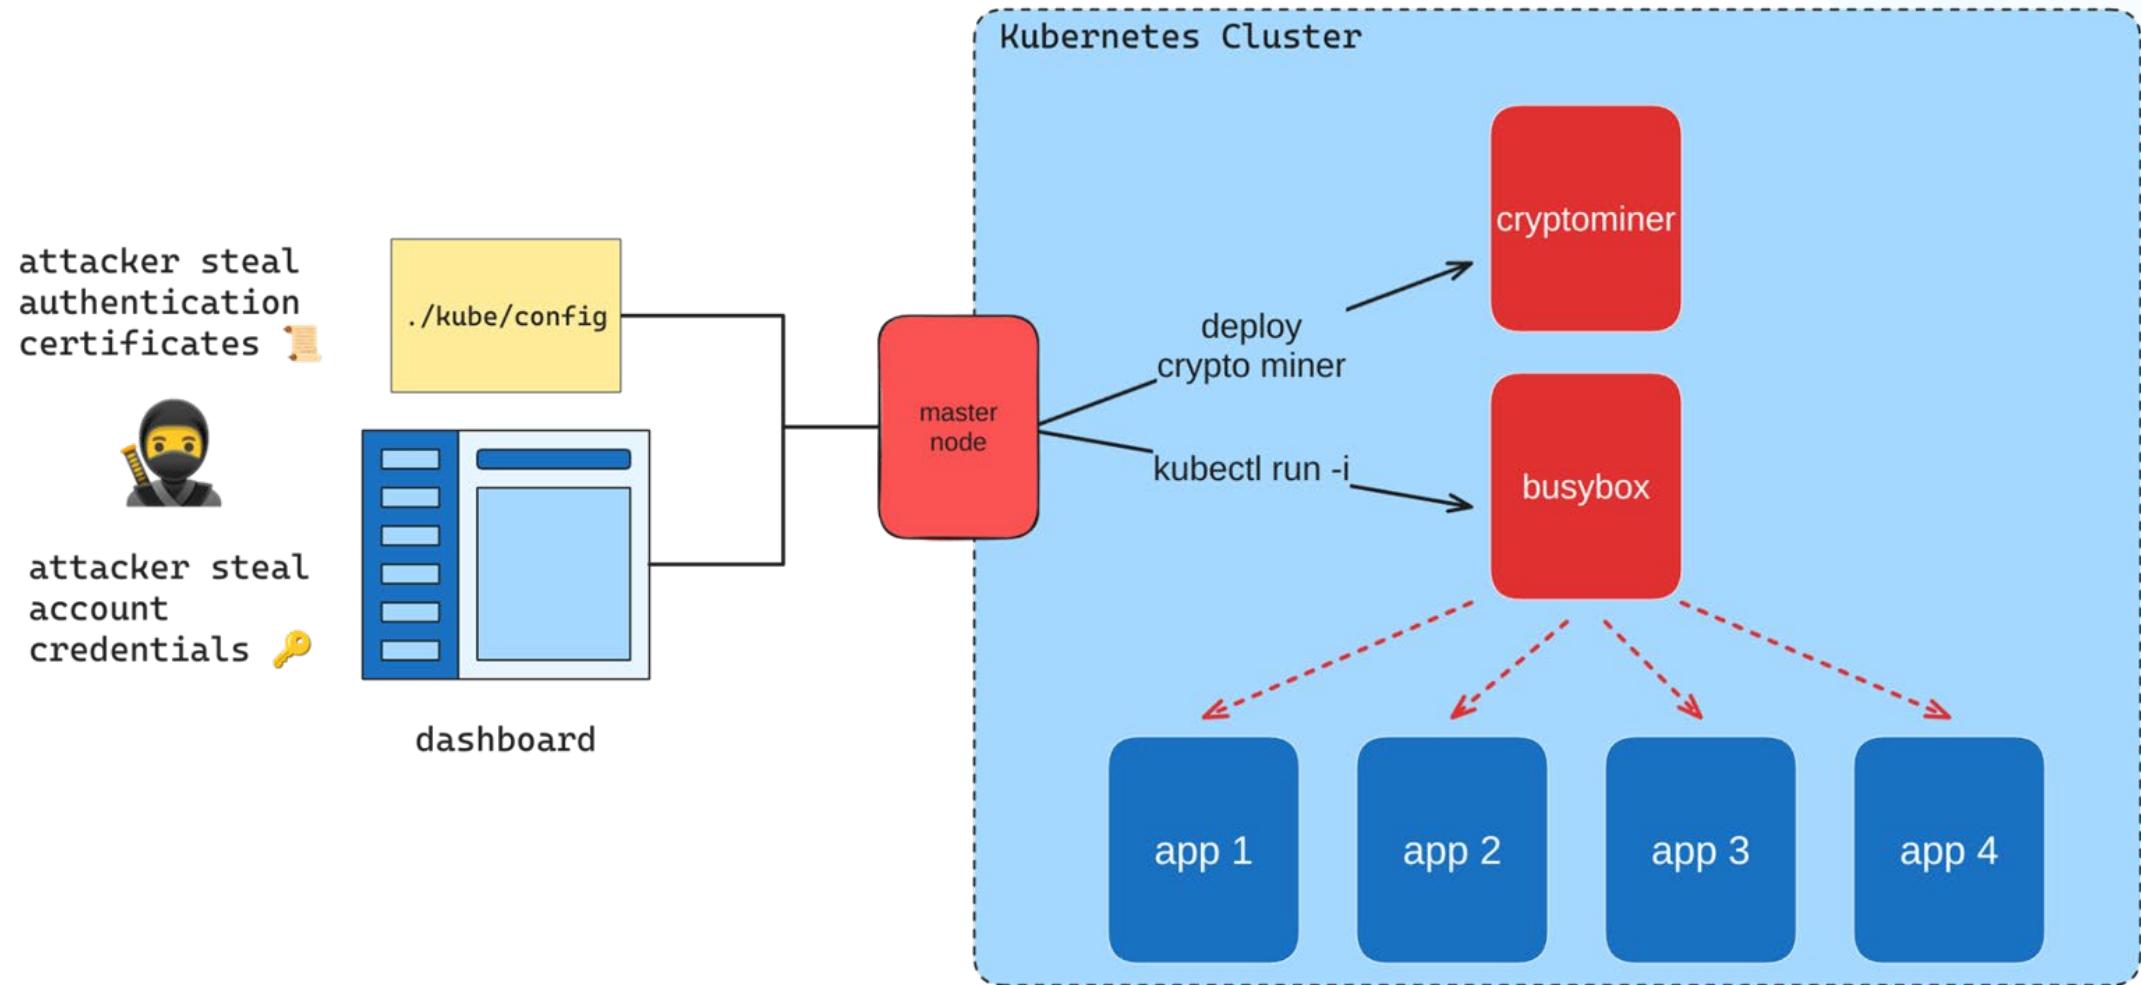
Execution

- Application exploit (RCE)
 - Exec into container
 - New container
 - Sidecar Injection

Exec Into The Container



Deploying New Containers



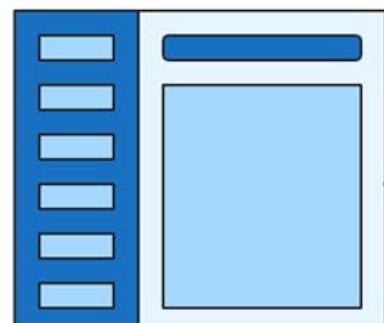
Sidecar Injection

attacker steal authentication certificates 📄



attacker steal account credentials 🔑

./kube/config



dashboard

Kubernetes Cluster

master node

kubectl patch

app 1

app 3

app 2

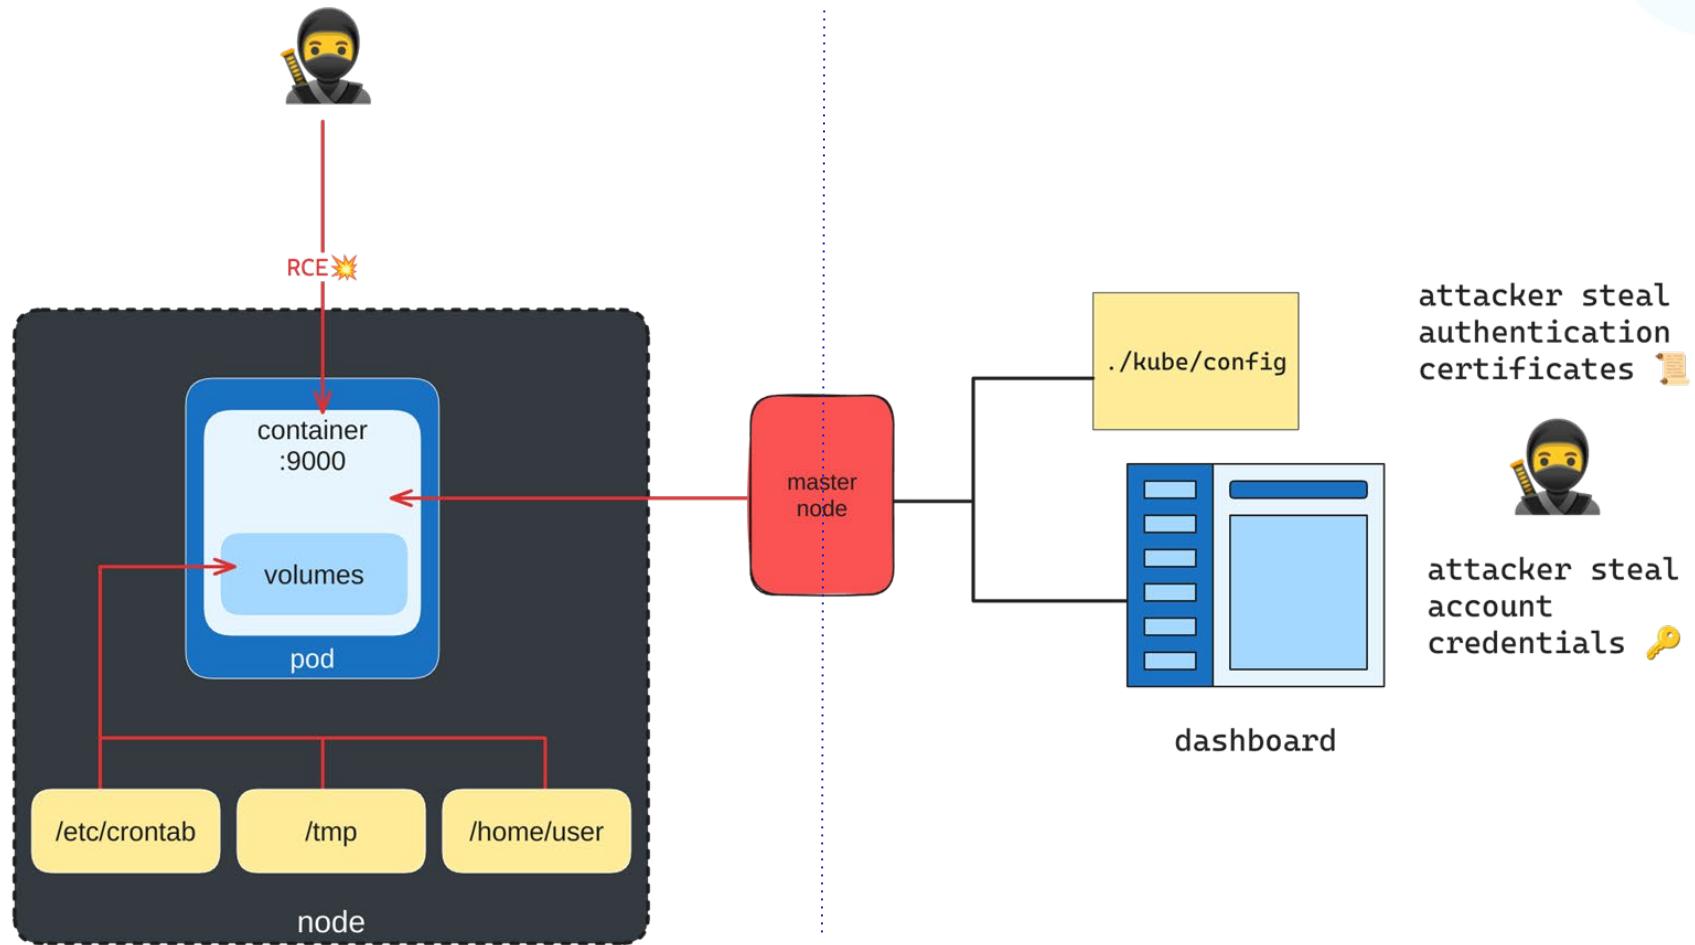
app 4

Persistence

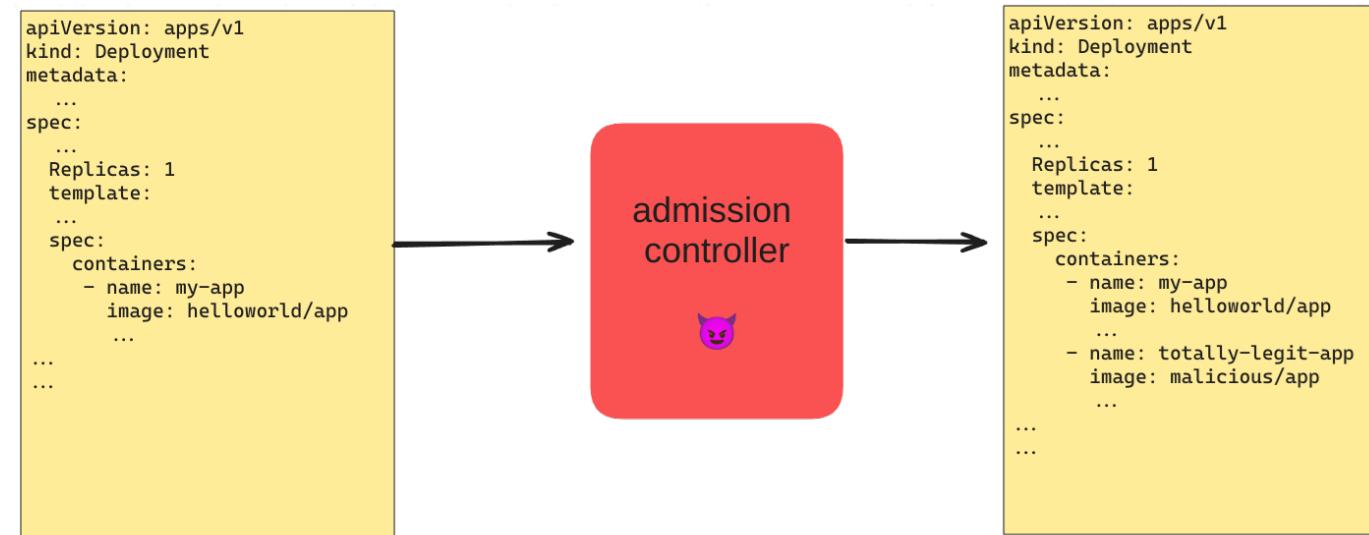
- Backdoor container
- Kubernetes CronJob
- Writable hostPath mount
- Malicious admission controller



Writable Hostpath Mount



Malicious Admission Controller

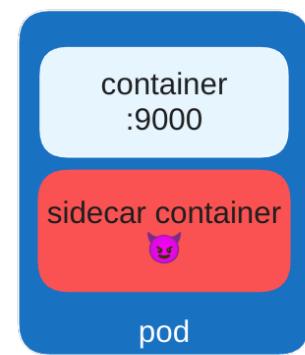


Legitimate user
deploying their
workload



Attacker controlling
admission controller

- malicious webhooks
- malicious mutation
- malicious validation
- tampered image



Privilege Escalation

- HostPath mount
- Access cloud resources
- Privileged container
- Cluster-admin binding



Privileged Container



Duffie Cooley



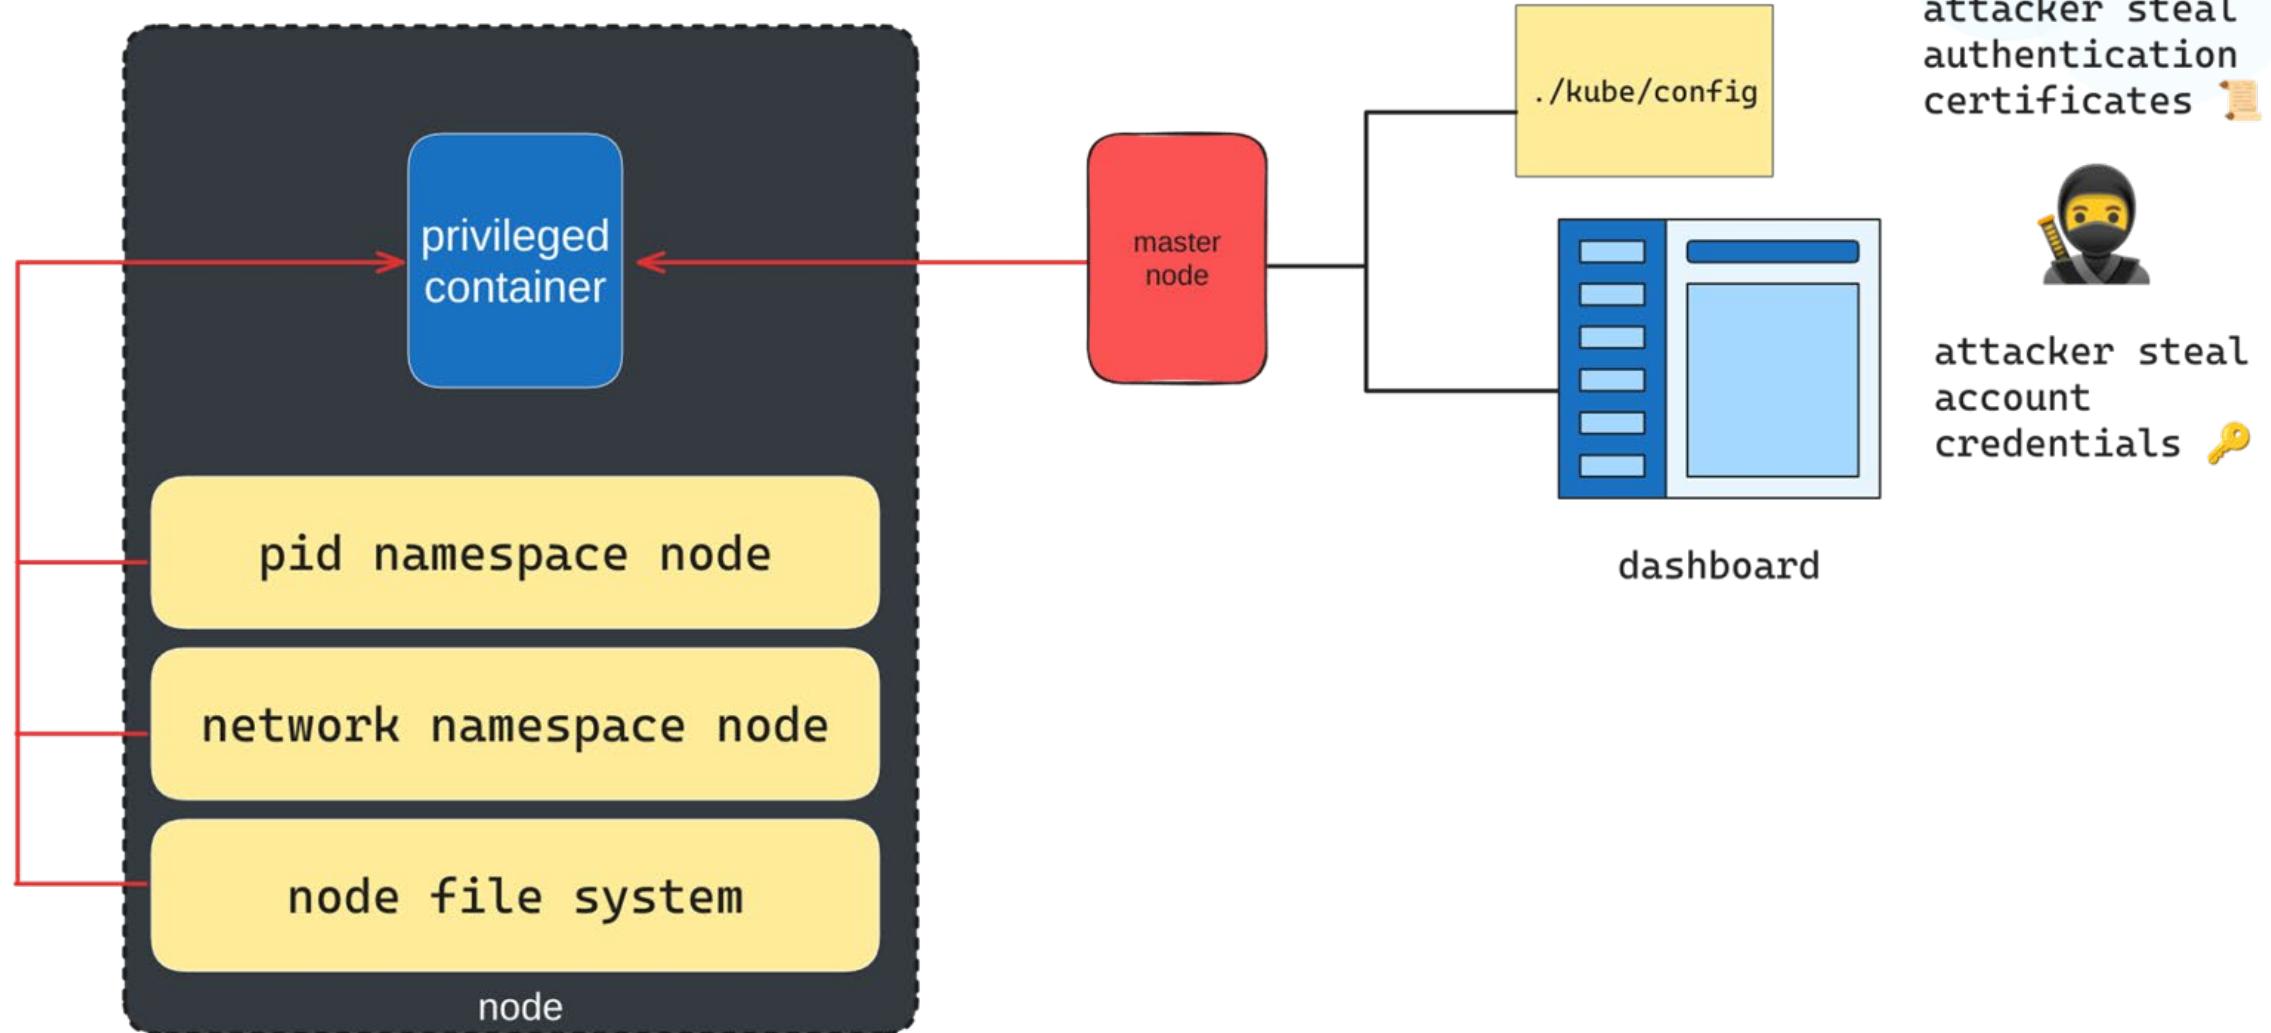
@maulion

...

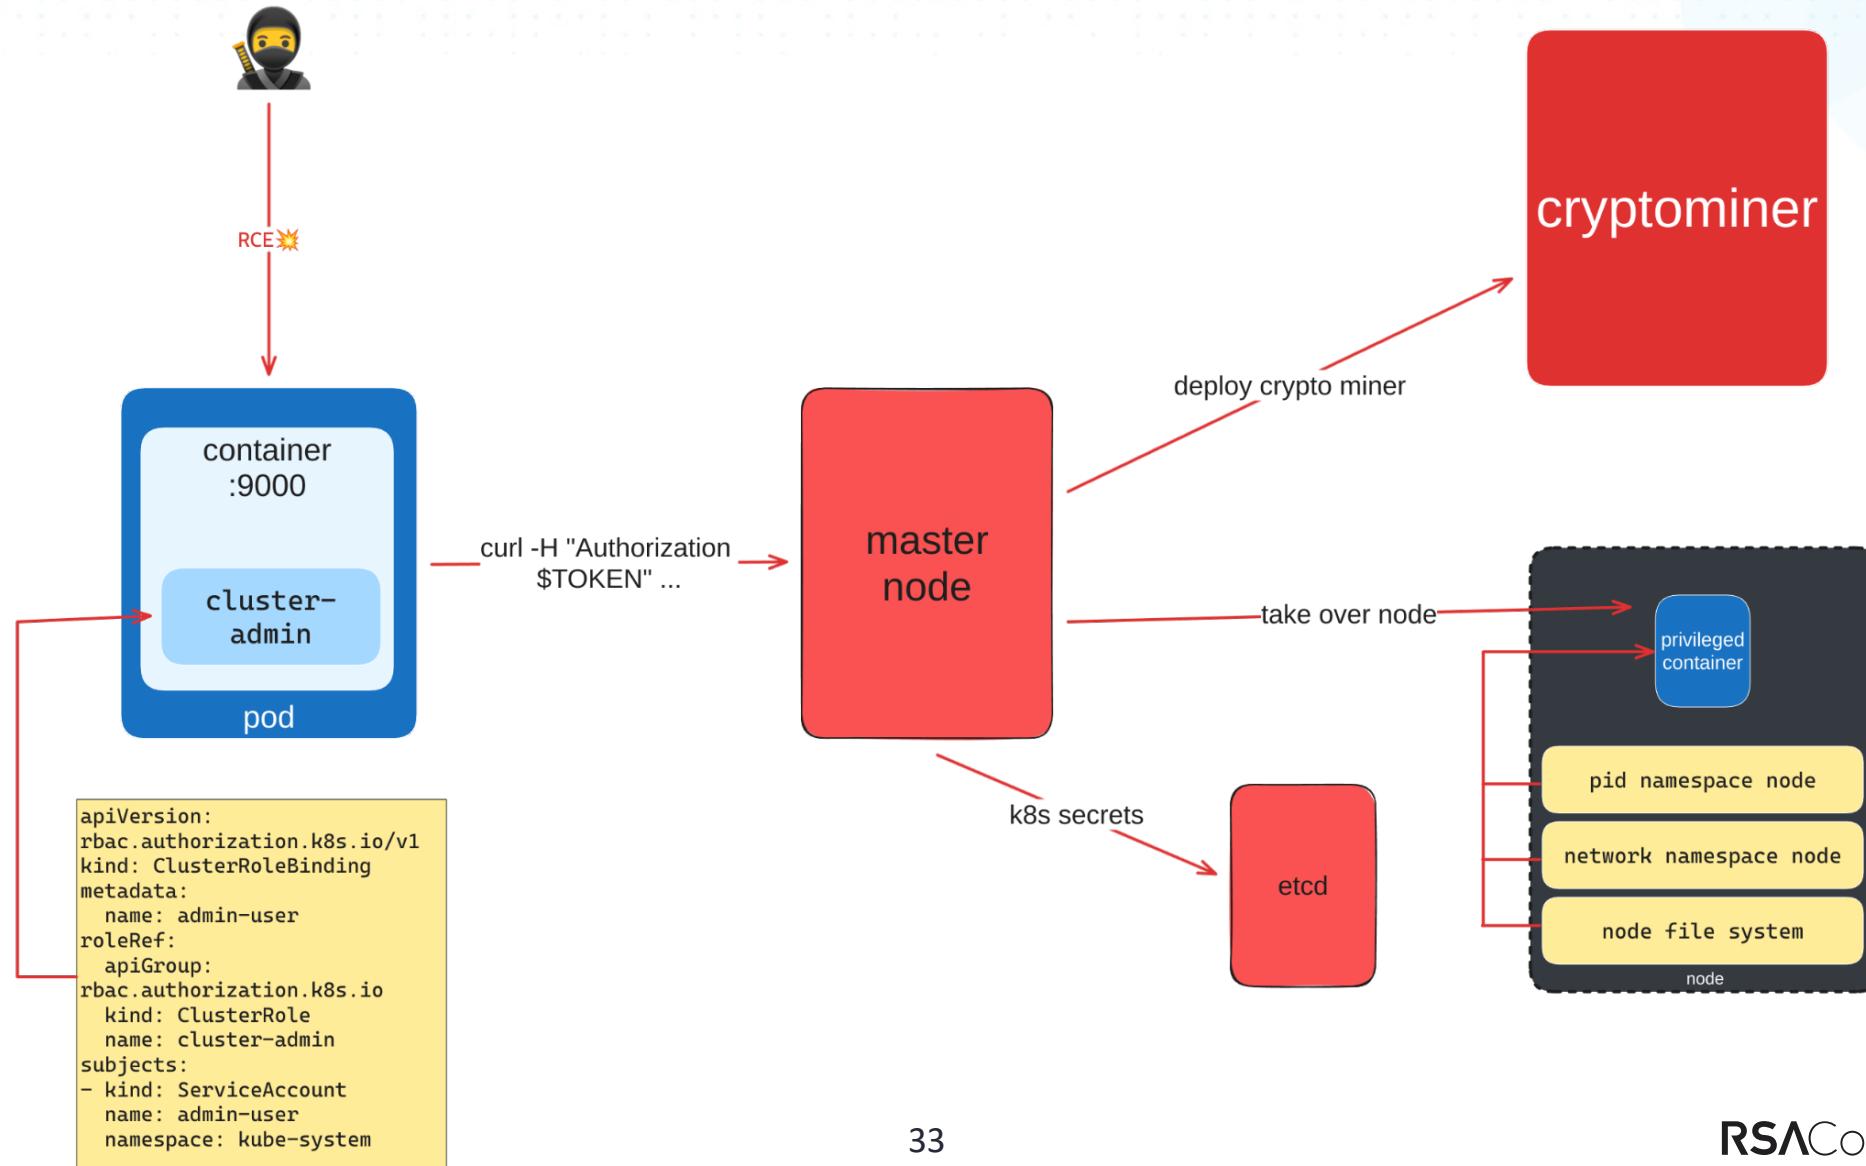
```
kubectl run r00t --restart=Never -ti --rm --image lol --  
overrides '[{"spec": {"hostPID": true, "containers":  
[{"name": "1", "image": "alpine", "command": ["nsenter", "-  
-mount=/proc/1/ns/mnt", "--", "/bin/bash"], "stdin":  
true, "tty": true, "securityContext": {"privileged": true}}]}]'
```

12:27 PM · May 17, 2019 · Twitter Web Client

A Container That Doesn't Contain Anything



Cluster-Admin Binding



Defense Evasion

- Clear container logs
- Delete Kubernetes events
- Pod / Container name similarity
- Connect from Proxy server



Credential Access

- List Kubernetes secrets
- Mount Service Principal
- Access container service account
- Applications credentials in configuration files
- Access managed identity credential
- Malicious admission controller



Discovery

- Access the Kubernetes API server
- Access Kubelet API
- Network mapping
- Access Kubernetes dashboard
- Instance Metadata API



Lateral Movement

- Access cloud resources
- Container service account
- Cluster internal networking
- Applications credentials in configuration files
- Writable volume mounts on the host
- CoreDNS poisoning



Collection

- Images from private registry



Impact

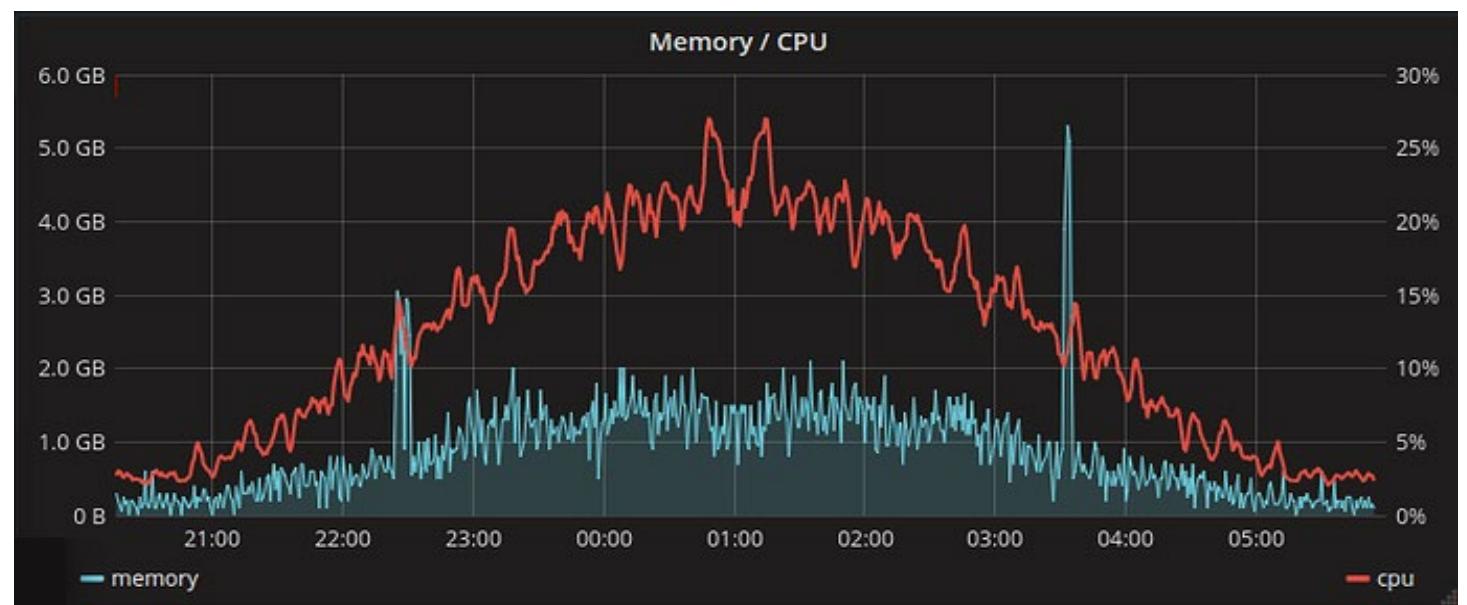
- Data destruction
 - Deleting PV, logs, events, etc
- Resource Hijacking
 - Deploying cryptominers, torrents, etc
- Denial of service
 - Deleting pods, deployments, statefulsets, services or ingress rules



Kubernetes Build-In Defenses

Resource Quotas and Limits

- Limits
 - CPU, Memory, etc.
- Requests
 - CPU, Memory, etc.

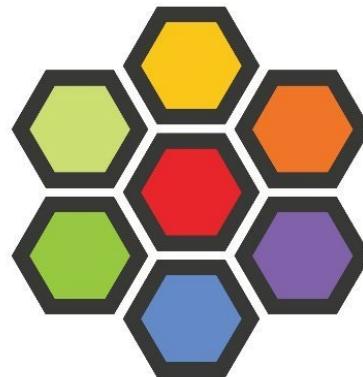


SecurityContext For Pods

- Privilege and Access Control
 - allowPrivilegeEscalation, privileged, runAsNonRoot, runAsUser, runAsGroup
- Capabilities Management
 - Capabilities: add, drop
- Filesystem and Mount Settings
 - readOnlyRootFilesystem
- Security Profiles
 - seLinuxOptions, seccompProfile, windowsOptions

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: security-context-demo
5 spec:
6   securityContext:
7     runAsUser: 1000
8     runAsGroup: 3000
9     fsGroup: 2000
10  volumes:
11    - name: sec-ctx-vol
12      emptyDir: {}
13  containers:
14    - name: sec-ctx-demo
15      image: busybox
16      command: [ "sh", "-c", "sleep 1h" ]
17      volumeMounts:
18        - name: sec-ctx-vol
19          mountPath: /data/demo
20      securityContext:
21        allowPrivilegeEscalation: false
```

Network Policies

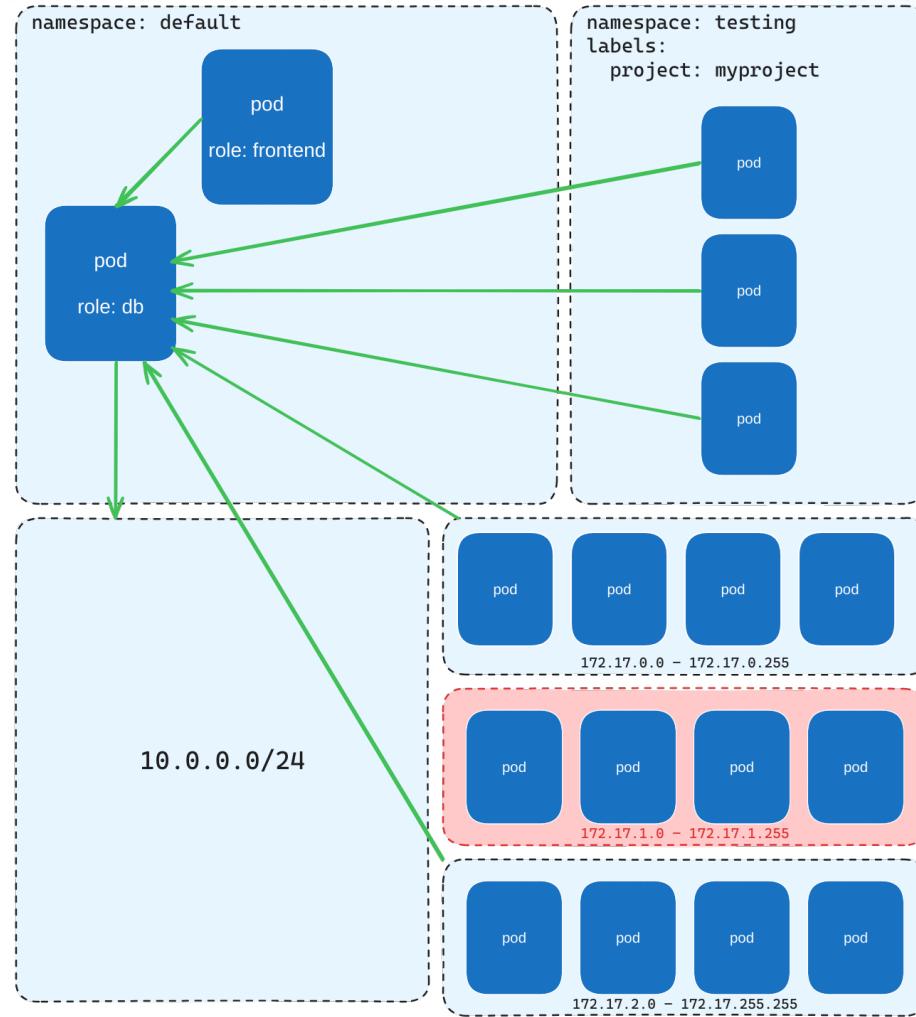


Network Policies

```

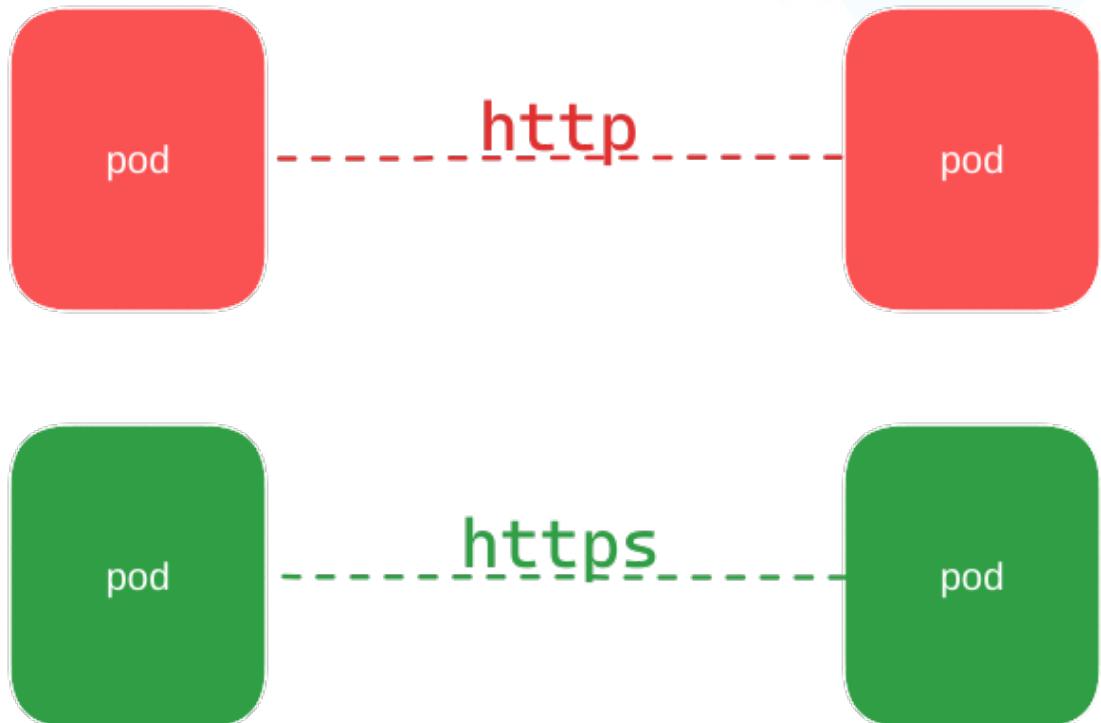
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
            except:
              - 172.17.1.0/24
        - namespaceSelector:
            matchLabels:
              project: myproject
        - podSelector:
            matchLabels:
              role: frontend
  ports:
    - protocol: TCP
      port: 6379
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
  ports:
    - protocol: TCP
      port: 5978

```



Encryption In Transit

- Kubernetes Certificate Authority
 - certificates.k8s.io
- Cert-manager.io
- Service mesh
 - Istio, Linkerd, Consul, etc



Encryption At Rest With EncryptionConfiguration

- Enable at-rest encryption for:
 - secrets, configmaps, events, *.apps, *.* , etc.
- Supports:
 - Identity (plain text), AESCBC, AESGCM
 - External KMS

```
10 apiVersion: apiserver.config.k8s.io/v1
 9 kind: EncryptionConfiguration
 8 resources:
 7   - resources:
 6     - secrets
 5     providers:
 4       - aescbc:
 3         keys:
 2           - name: key1
 1             secret: 6YbX+1+UK9ym6BP2tLVi0JqYYZguVgzgMnrew7oK
 0   - identity: {}
 1
```

```
<option-configuration.yaml | unix | utf-8 | yaml | 91% | 11:1
<-encryption/encryption-configuration.yaml" [noeol] 12L, 259B
```

Pod Security Policies (PSPs)

- Deprecated in Kubernetes v1.21, and removed from Kubernetes in v1.25.

Role

RoleBinding

PodSecurityPolicy

ClusterRole

ClusterRole Binding

Protect Cluster Against Privilege Pods

PodSecurityPolicy

```

1 apiVersion: policy/v1beta1
2 kind: PodSecurityPolicy
3 metadata:
4   name: restrictive ←
5 spec:
6   privileged: false
7   hostNetwork: false
8   allowPrivilegeEscalation: false
9   defaultAllowPrivilegeEscalation: false
10  hostPID: false
11  hostIPC: false
12  runAsUser:
13    rule: RunAsAny
14  fsGroup:
15    rule: RunAsAny
16  seLinux:
17    rule: RunAsAny
18  supplementalGroups:
19    rule: RunAsAny
20  volumes:
21  - 'configMap'
22  - 'downwardAPI'
23  - 'emptyDir'
24  - 'persistentVolumeClaim'
25  - 'secret'
26  - 'projected'
```

ClusterRole

```

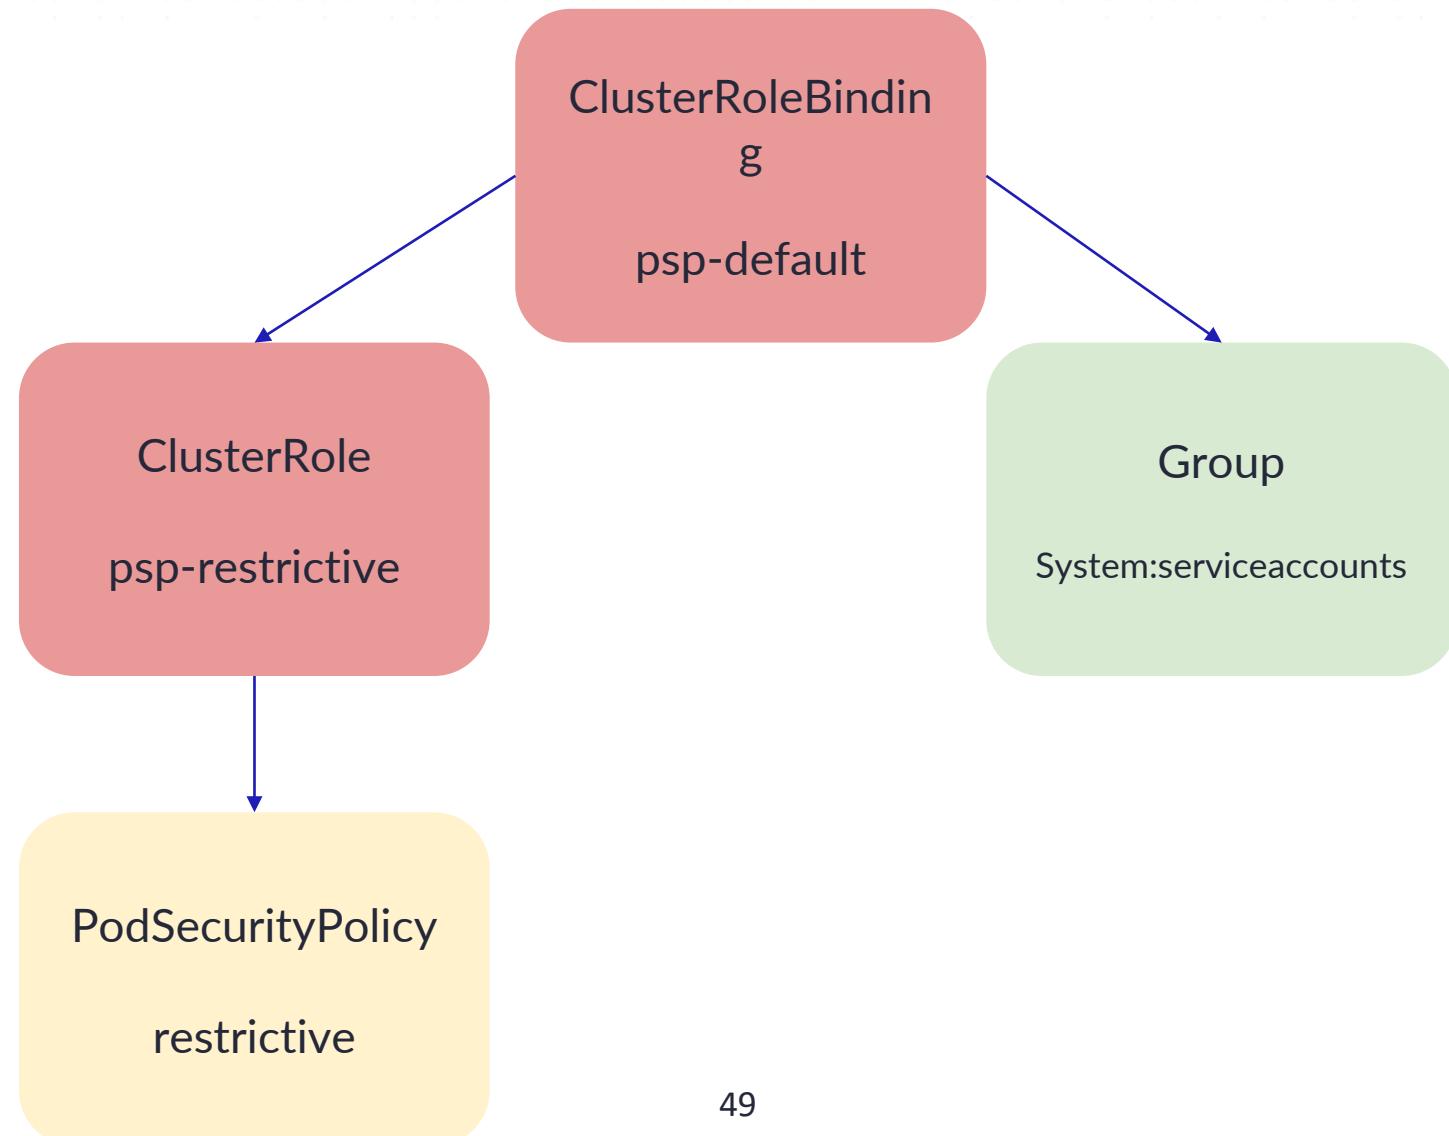
1 kind: ClusterRole
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: psp-restrictive
5 rules:
6   - apiGroups:
7     - extensions
8     resources:
9       - podsecuritypolicies
10    resourceNames:
11      - restrictive ←
12    verbs:
13      - use
```

ClusterRoleBinding

```

1 kind: ClusterRoleBinding
2 apiVersion: rbac.authorization.k8s.io/v1
3 metadata:
4   name: psp-default
5 subjects:
6   - kind: Group
7     name: system:serviceaccounts ←
8     namespace: kube-system
9 roleRef:
10  kind: ClusterRole ←
11  name: psp-restrictive ←
12  apiGroup: rbac.authorization.k8s.io
```

Protect Cluster Against Privilege Pods



Protect Cluster Against Privilege Pods

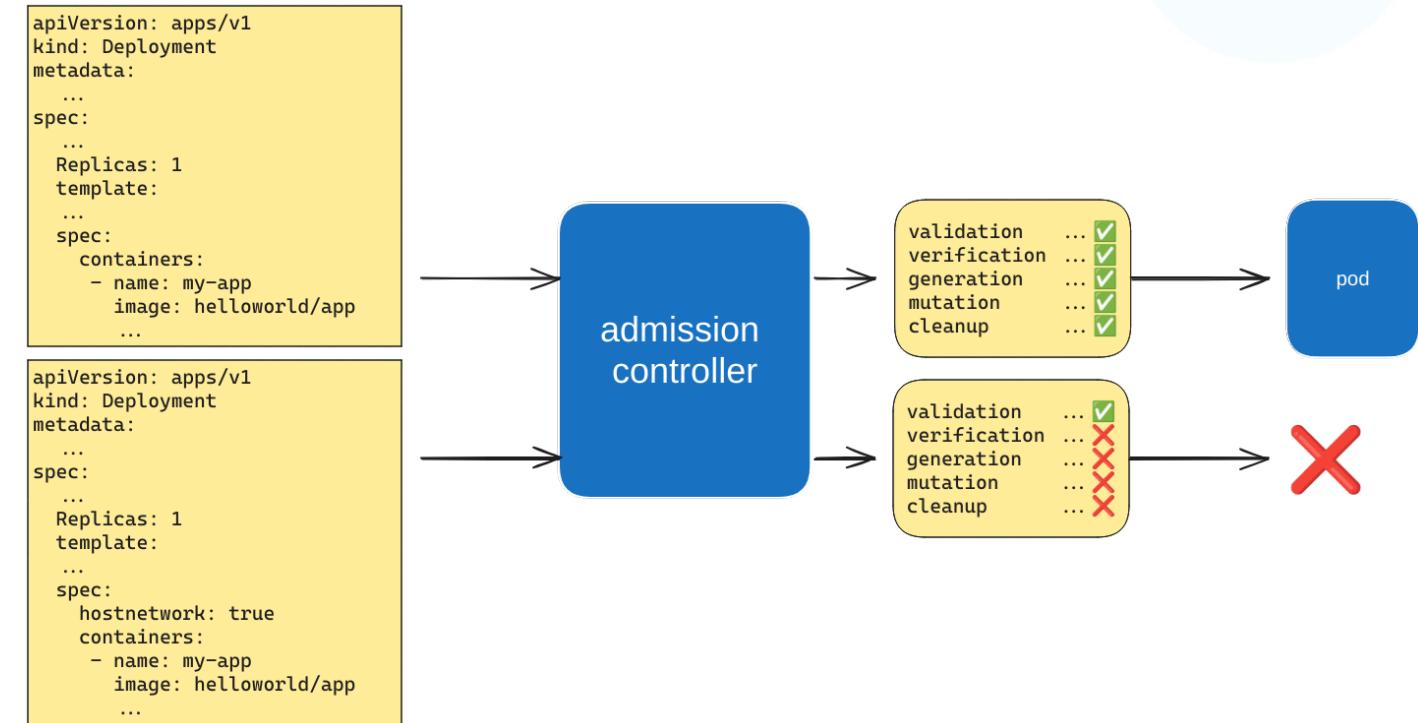
```
2 kind: Deployment
3 metadata:
4   name: nginx-hostnetwork-deployment
5   namespace: default
6   labels:
7     app: nginx
8 spec:
9   replicas: 1
10  selector:
11    matchLabels:
12      app: nginx
13  template:
14    metadata:
15      labels:
16        app: nginx
17    spec:
18      containers:
19        - name: nginx
20          image: nginx:1.15.4
21          hostNetwork: true
```

Controlled By: Deployment/nginx-hostnetwork-deployment
Replicas: 0 current / 1 desired
Pods Status: 0 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
 Labels: app=nginx
 pod-template-hash=597c4cff45
 Containers:
 nginx:
 Image: nginx:1.15.4
 Port: <none>
 Host Port: <none>
 Environment: <none>
 Mounts: <none>
 Volumes: <none>
 Conditions:
 Type Status Reason

 ReplicaFailure True FailedCreate
Events:
 Type Reason Age From Message
 Warning FailedCreate 13s (x13 over 33s) replicaset-controller Error creating: pods "nginx-hostnetwork-deployment-597c4cff45-" is forbidden: unable to validate against any pod security policy: [spec.securityContext.hostNetwork: Invalid value: true: Host network is not allowed to be used]

Pod Security Admission

- Admission Controller
- Validation, mutation, cleanup
Kubernetes resources
- OPA GateKeeper, Kyverno,
jsPolicy, etc



Audit Log

- When to Log:
 - RequestReceived, ResponseStarted, ResponseComplete, Panic
- What to Log:
 - None, Metadata, Request, RequestResponse
- Where to Log:
 - Filesystem, Webhook

```
15 apiVersion: audit.k8s.io/v1 # This is required.
14 kind: Policy
13 # Don't generate audit events for all requests in RequestR
12 omitStages:
11 - "RequestReceived"
10 rules:
  9 # Log pod changes at Request level
  8 - level: Request
    7 resources:
      6 - group: ""
        5 resources: ["pods"]
  4 # Log pod changes at RequestResponse level
  3 - level: RequestResponse
    2 resources:
      1 - group: ""
        0 resources: ["pods"]
```

~
~
~
~

NORMAL audit-policy.yaml unix | utf-8 | yaml 100% 16:2

Kubernetes Cluster Hardening

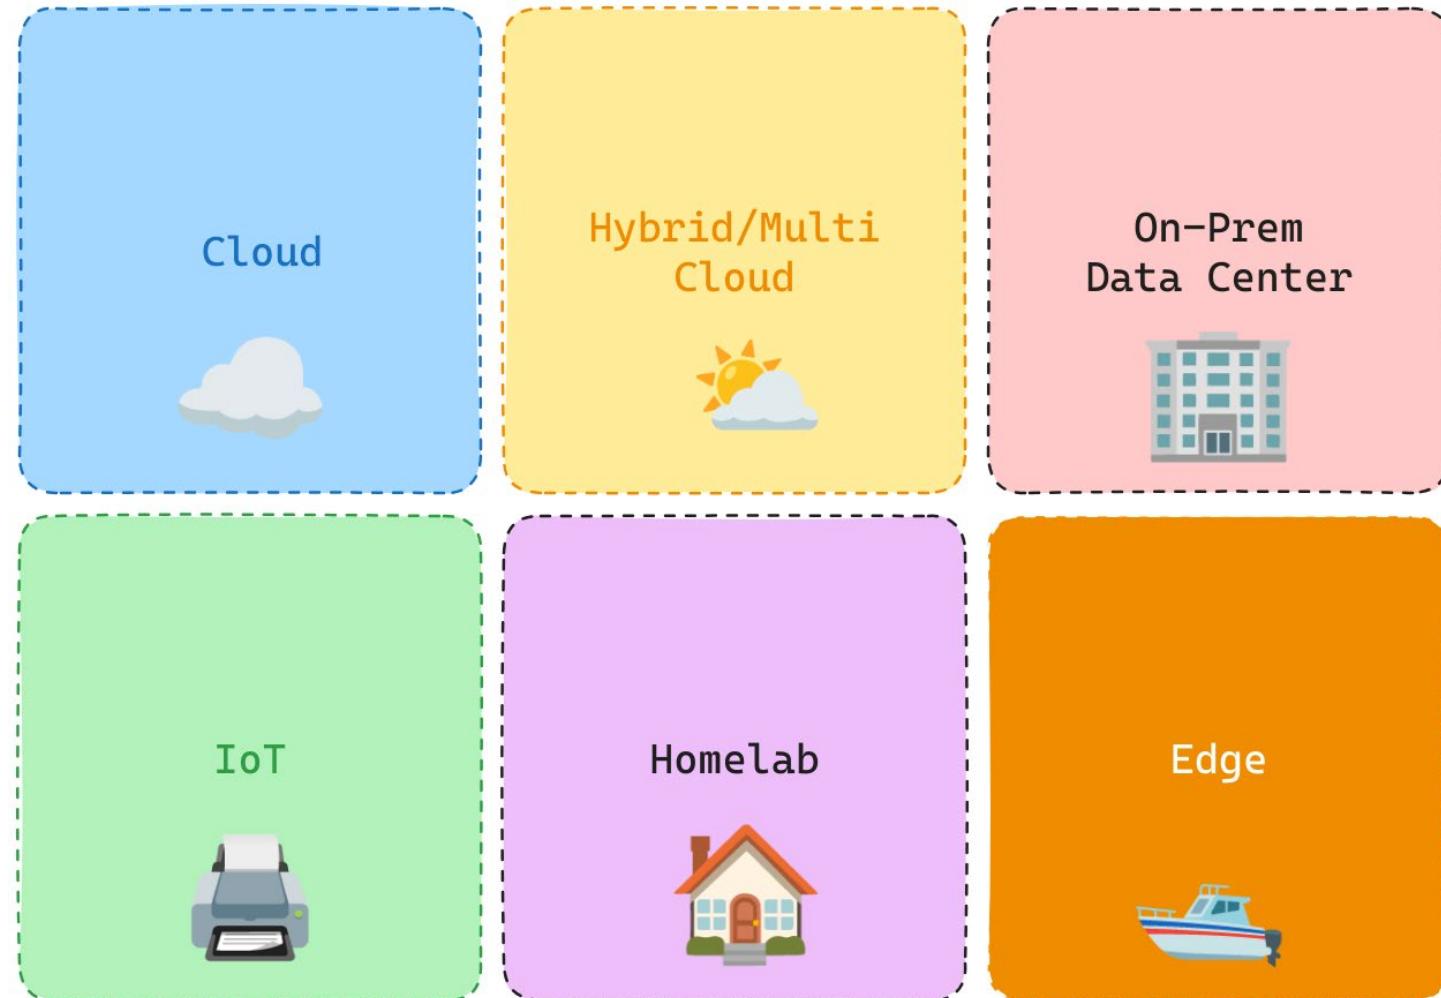
- Isolate Kubernetes Nodes
 - Don't expose nodes directly, use ingress controller, separate control and data flow
- Protect Kubelet
 - Disable anonymous authentication
 - Enable mTLS authentication on the kubelet's HTTPS endpoint
- Protect etcd with TLS, Firewall and Encryption
- Use Third-Party Authentication for API Server



What's Next?

Kubernetes Evolution (Edge & Hybrid Cloud)

- Kubeedge
- SuperEdge
- Akri
- OpenYurt
- K3s
- Microk8s
- Minikube
- ...



Next Time You Deploy Kubernetes – Part 1

- Don't deploy images from untrusted sources in your cluster, regularly scan images for vulnerabilities
- Don't store sensitive data on kubernetes secrets, instead use an external KMS to manage secrets or enable EncryptionConfiguration
- Encryption in transit, generate TLS certificates for your workloads using the certificates.k8s.io api or use a third party solution like cert-manager.io

Next Time You Deploy Kubernetes – Part 2

- Enforce SecurityContext and limit resource consumption for your workloads
- Network Segmentation, choose CNI plugin that supports network policies to limit communication between your services and monitor network traffic
- Configure Role-Based Access Control (RBAC) policies via an Admission Controller or PSPs always following the principle of least privilege. Never use cluster-admin in your workloads

Next Time You Deploy Kubernetes – Part 3

- Turn on audit logging and monitor unusual or unwanted api calls, e.g. authentication failures.
- Keep Kubernetes version up to date, update and patch regularly
- Hardening control plane components and lock down kubelet
- Adopt GitOps practices for managing your Kubernetes clusters to enhance security, auditability, and consistency through declarative configuration, automation, and tight integration with version control systems

Thanks